

STATS-413 HW-1

Author: Hochan Son

UID: 206547205

Date: @October 10, 2025

Problem 1: For linear and logistic regressions,

assuming $s_i = \sum_j \beta_j x_{ij} = x_i^T \beta$.

(1) Find the log-likelihood function

$$J(\beta) = \sum_{i=1}^n \log p(y_i | s_i = x_i^T \beta)$$

For a single observation, the likelihood is, where σ = sigmoid function:

$$p(y_i | s_i) = \sigma(s_i) = \frac{e^{y_i s_i}}{1 + e^{s_i}}$$

Taking the log, for the chain rule:

$$\log p(y_i|s_i) = y_i s_i - \log(1 + e^{s_i})$$

for $s_i = x_i^T \beta$,

$$\log p(y_i|s_i = x_i^T \beta) = y_i x_i^T \beta - \log(1 + e^{x_i^T \beta})$$

MLE (maximum log likelihood):

$$\sum_{i=1}^n \log p(y_i|s_i = x_i^T \beta) = \sum_{i=1}^n y_i x_i^T \beta - \log(1 + e^{x_i^T \beta}) \quad (1)$$

The log likelihood function:

$$J(\beta) = \sum_{i=1}^n [\log p(y_i|s_i = x_i^T \beta)] = \sum [y_i s_i - \log(1 + e^{s_i})] \quad (2)$$

Plug $s_i = x_i^T \beta$ into (2),

$$J(\beta) = \sum_{i=1}^n \log p(y_i|s_i = x_i^T \beta) = \sum_{i=1}^n y_i x_i^T \beta - \log(1 + e^{x_i^T \beta}) \quad (3)$$

Therefore, (1) == (3)

(2) Find the gradient $J'(\beta)$ by calculating $\frac{dJ}{d\beta}$, and write down the gradient ascent algorithm.

$$\begin{aligned} J'(\beta) &= \frac{d}{d\beta} \sum_{i=1}^n \log p(y_i|s_i = x_i^T \beta) \\ &= \frac{d}{d\beta} \sum_{i=1}^n y_i x_i^T \beta - \log(1 + e^{x_i^T \beta}), \\ &\text{where } y_i = \text{const} \end{aligned}$$

$$= y_i \frac{d}{d\beta} \sum_{i=1}^n x_i^T \beta - \frac{dJ}{d\beta} \log(1 + e^{x_i^T \beta}) \quad (4)$$

Let's substitute (1) : $u = x_i^T \beta$, (2) : $v = e^u$, (3) : $w = \log(1 + v)$, chain rule:

$$\begin{aligned} & \frac{d}{d\beta} \log(1 + e^{x_i^T \beta}) \\ &= \frac{dw}{dv} * \frac{dv}{du} * \frac{du}{d\beta} \end{aligned}$$

For (1),

$$\frac{du}{d\beta} x_i^T \beta = \frac{du}{d\beta} (x_{i1}\beta + x_{i2}\beta + \dots + x_{ip}\beta) = x_i \text{ (scalar)}$$

For (2),

$$\begin{aligned} \frac{dv}{du} v &= \frac{dv}{du} e^u = e^u \\ &= e^{x_i^T \beta} \end{aligned}$$

For (3),

$$\begin{aligned} \frac{dw}{dv} \log(1 + v) &= \frac{1}{1 + v} \times \frac{dw}{dv} (1 + v) = \frac{1}{1 + v} \times 1 = \frac{1}{1 + v} \\ &\text{since } v = e^{x_i}, \\ &= \frac{1}{1 + e^{x_i^T \beta}} \end{aligned}$$

To combine (1) (2), and (3),

$$\begin{aligned} (1)(2)(3) &= \frac{1}{1 + e^{x_i^T \beta}} \times e^{x_i^T \beta} \times x_i \\ &= \frac{e^{x_i^T \beta}}{1 + e^{x_i^T \beta}} \times x_i \end{aligned}$$

Let's substitute $x_i^T \beta$ to α :

$$= \frac{e^\alpha}{1 + e^\alpha} x_i = \text{sigmoid function}(\sigma) \\ = \sigma(\alpha)x_i$$

$$\text{ans} = \sigma(x_i^T \beta) x_i \\ (\text{since } s_i = x_i^T \beta), \\ = s_i x_i$$

Problem 2: For logistic regression based on a simple neural network, with

$$s_i = \beta_0 + \sum_{k=1}^d \beta_k h_{ik}, \quad (5)$$

$$h_{ik} = \max(0, s_{ik}), -ReLU \quad (6)$$

$$s_{ik} = \alpha_{k0} + \sum_{j=1}^p \alpha_{kj} x_{ij} \quad (7)$$

(1) Using one-dimensional and two-dimensional input x_i , explain how the above neural network defines a piecewise linear function $s_i = f(x_i)$

For (5) is a linear function, if h_{ik} is linear,

$$s_i = \beta_0 + \beta_1 h_{i1} + \beta_2 h_{i2} + \dots + \beta_d h_{id}$$

However, for (6), given h_{ik} is ReLU (Rectified Linear Unit), it is conditional linear (or a.k.a. piecewise linear)

When $s_{ik} > 0$, s_{ik} is linear, otherwise 0

Given that the conditional linearity makes it work like a piecewise linear function with respect to $s_i = f(x_i)$

(2) Explain how the last line can be expressed in terms of vectors and matrices.

For (7),

$$\begin{aligned}s_{ik} &= \alpha_{k0} + \sum_{j=1}^p \alpha_{kj} x_{ij} \\ &= a_{k0} + a_{k1}x_{i1} + a_{k2}x_{i2} + \dots + a_{kp}x_{ip}\end{aligned}$$

It can be simplified to

Let $\alpha_0 = [\alpha_0 \quad \alpha_1 \quad \alpha_2 \quad \dots \quad \alpha_p]$; ($p \times 1$ matrix)

Let $A^T = \begin{bmatrix} a_{00} & a_1 & a_2 & \dots & a_p \\ a_{01} & & & & \end{bmatrix}$; ($d \times p$ matrix)

Let $X = \begin{bmatrix} X_0 \\ X_1 \\ X_2 \\ \dots \\ X_p \end{bmatrix}$, ($1 \times p$ matrix)

$$s_{ik} = \alpha_{k0} + \sum_{j=1}^p \alpha_{kj} x_{ij} = \alpha_{k0} + \alpha_k^T x_i = \alpha_0 + A^T X$$

(3) Let $\theta = (\beta_k, \alpha_{kj}, \forall_k, j)$. Find the gradient $J'(\theta)$ by calculating $\frac{\partial J}{\partial \beta_k}$ and $\frac{\partial J}{\partial \alpha_{kj}}$, and write down the gradient ascent algorithm with momentum.

1. Calculate $\frac{\partial J}{\partial \beta_k}$, using chain rule:

- to find $\frac{\partial J}{\partial s_i}$, where p_i is the predicted probability

$$\frac{\partial J}{\partial \beta_k} = \sum_{i=1}^n \frac{\partial J}{\partial s_i} \times \frac{\partial S_i}{\partial \beta_k}$$

$$\frac{\partial J}{\partial s_i} = y_i - \frac{e_i^s}{1 + e_i^s} = y_i - \sigma(s_i) = y_i - p_i$$

- to find $\frac{\partial s_i}{\partial \beta_k}$:

$$\frac{\partial s_i}{\partial \beta_k} = h_{ik}$$

therefore,

$$\frac{\partial J}{\partial \beta_k} = \sum_{i=1}^n (y_i - p_i) h_{ik}$$

- For h_{ik} (= ReLU),

$$h_{ik} = \max\{0, \alpha_{k0} + \sum_{j=1}^p \alpha_{kj} x_{ij}\}$$

- For β_0 ,

$$\frac{\partial J}{\partial \beta_k} = \sum_i (y_i - p_i)$$

- For B_k ,

$$\frac{\partial J}{\partial \beta_k} = \sum_{i=1}^n (y_i - p_i) \max\{0, \alpha_{k0} + \sum_{j=1}^p \alpha_{kj} x_{ij}\}$$

2. Calculate $\frac{\partial J}{\partial \alpha_{kj}}$, using chain rule:

$$\begin{aligned} \frac{\partial J}{\partial \alpha_{kj}} &= \frac{\partial J}{\partial s_i} \frac{\partial s_i}{\partial h_{ik}} \frac{\partial h_{ik}}{\partial \alpha_{kj}} = \frac{\partial J}{\partial s_i} \frac{\partial s_i}{\partial \alpha_{kj}} \\ &= \sum_i (y_i - \frac{e_i^s}{1 + e_i^s}) (\beta_i \times \max\{0, x_{ij}\}) \end{aligned}$$

3. For $J'(\theta)$,

- - - - -

$$J'(\theta) = \left[\frac{\partial J}{\partial \beta_1}, \dots, \frac{\partial J}{\partial \beta_d}, \frac{\partial J}{\partial \alpha_1}, \dots, \frac{\partial J}{\partial \alpha_{dp}} \right]^T$$

4. Gradient Ascent algorithm with momentum

$$\begin{aligned}\theta^{(t+1)} &= \theta^{(t)} + \mu_t V^{(t)} \\ V^{(t)} &= \gamma v^{t-1} + J'(\theta)\end{aligned}$$

Problem 3: Write Python code to implement the algorithm in Problem 2, where $x_i = (x_{i,1}, x_{i,2})$ Uniform[0, 1]², i.e., a unit square, and $y_i = 1(x_{i,1}^2 + x_{i,2}^2 < 1)$, i.e., within a unit circle. Plot the decision boundary defined by the learned neural network.


I will introduce three Neural network models and compare the performance benchmark with data pipeline diagram

1. NeuralNetwork-Model 1: Adam Optimizer with Full batch load
2. NeuralNetwork-Model 2: Adam Optimizer with Fully Optimized Python code and mini batch
3. NeuralNetwork-Model3: Stochastic SGD with Momentum


Overview of the Data pipeline Algorithms

- Model Selection Layer: NN-Adam, NN-AdamFullOptimized, and NN-SGD+Momentum
- Training pipeline with Epoch based
- Prediction layer: Predict-probability and Predict
- Source Code: I've attached as distinct files.


- hw1.py: Main function
- NeuralNetworkAdam.py : Adam optimizer with full batch
- NeuralNetworkAdamFullyOptimized.py
- Stochastic Gradient Descent with Momentum
- Boundary detection with three different models and optimizer performance comparison
 - As seen the graphs below, SGD+Momentum has absolutely great accuracy (Val Acc:0.9984) by outperformed other Adam Optimizer models




1. Full Neural Network Implementation Diagram




2. Neural Network Layer-by-Layer Flow




3. Forward pass




4. Backward Pass (Backpropagation)




5. Adam Optimizer Update Process



6. SGD with Momentum Update Process



7. Data flow entire system



Codebase explanation

- Main function

```
"""
```

STATS 413 - Homework 1

Neural Network Implementation Comparison

Compares three optimization algorithms:

1. Adam (full-batch)
2. Adam (mini-batch with optimizations)
3. SGD with Momentum (mini-batch)

Task: Learn circular decision boundary ($x_1^2 + x_2^2 < 1$)

```
"""
```

```
import numpy as np
import matplotlib.pyplot as plt
import time
from NeuralNetworkAdam import NeuralNetworkAdam
from NeuralNetworkAdamFullyOptimized import NeuralNetworkAdamFullyOptimized
from NeuralNetworkSGDMomentum import NeuralNetworkSGDMomentum
```

```
def generate_data(n_samples):
```

```
    """
```

Generate training data: uniform in $[0,1]^2$, labeled by unit circle

Args:

n_samples: Number of samples to generate

Returns:

X: Input features of shape (n_samples, 2)

y: Binary labels of shape (n_samples,)

```
"""
```

```
X = np.random.uniform(0, 1, size=(n_samples, 2))
```

```
y = (X[:, 0]**2 + X[:, 1]**2 < 1).astype(int)
```

```
return X, y
```

```

def plot_decision_boundary(ax, model, X_val, y_val, title, val_acc):
    """
    Plot decision boundary for a trained model

    Args:
        ax: Matplotlib axis object
        model: Trained neural network model
        X_val: Validation data
        y_val: Validation labels
        title: Plot title
        val_acc: Validation accuracy
    """

    # Create meshgrid for decision boundary
    x1_min, x1_max = -0.1, 1.1
    x2_min, x2_max = -0.1, 1.1
    xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max, 300),
                           np.linspace(x2_min, x2_max, 300))
    X_grid = np.c_[xx1.ravel(), xx2.ravel()]

    # Get model predictions
    Z = model.predict_proba(X_grid).reshape(xx1.shape)

    # Plot filled contour
    contourf = ax.contourf(xx1, xx2, Z, levels=20, cmap='RdYlBu', alpha=0.7)
    ax.contour(xx1, xx2, Z, levels=[0.5], colors='black', linewidths=3)

    # Sample validation points for visualization
    sample_size = min(500, len(y_val))
    sample_idx = np.random.choice(len(y_val), sample_size, replace=False)

    ax.scatter(X_val[sample_idx][y_val[sample_idx]==1, 0],
               X_val[sample_idx][y_val[sample_idx]==1, 1],
               c='blue', marker='o', edgecolors='k', s=15, alpha=0.6,
               label='Inside (y=1)')

```

```

        ax.scatter(X_val[sample_idx][y_val[sample_idx]==0, 0],
                    X_val[sample_idx][y_val[sample_idx]==0, 1],
                    c='red', marker='s', edgecolors='k', s=15, alpha=0.6,
                    label='Outside (y=0)')

# Plot true boundary
theta = np.linspace(0, np.pi/2, 100)
ax.plot(np.cos(theta), np.sin(theta), 'g--', linewidth=3, label='True boundary')

ax.set_xlabel('$x_1$', fontsize=12)
ax.set_ylabel('$x_2$', fontsize=12)
ax.set_title(f'{title}\n(Val Acc: {val_acc:.4f})', fontsize=12, fontweight='bold')
ax.legend(fontsize=9, loc='upper right')
ax.set_xlim(x1_min, x1_max)
ax.set_ylim(x2_min, x2_max)
ax.set_aspect('equal')
ax.grid(True, alpha=0.3)

return contourf

def main():
    """Main training and evaluation function"""

    # Set random seed for reproducibility
    np.random.seed(2024)

    print("=" * 80)
    print("STATS 413 - HW1: Neural Network Optimizer Comparison")
    print("Task: Learn quarter circle decision boundary")
    print("=" * 80)
    print()

    # =====
    # Generate Data

```

```

# =====
print("Generating data...")
N_train = 10000
N_val = 5000

X_train, y_train = generate_data(N_train)
X_val, y_val = generate_data(N_val)

print(f"Training samples: {N_train}")
print(f"Validation samples: {N_val}")
print(f"Positive samples (inside circle): {np.sum(y_train)}/{N_train} = {np.mean(y_train)[:2%]}")
print()

# =====
# Model 1: Adam (Full-Batch)
# =====
print("=" * 80)
print("MODEL 1: Adam Optimizer (Full-Batch)")
print("=" * 80)

model1 = NeuralNetworkAdam(
    input_dim=2,
    hidden_dim=128,
    learning_rate=0.001,
    beta1=0.9,
    beta2=0.999
)

print(f"Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid")
print(f"Optimizer: Adam (lr={model1.lr}, β₁={model1.beta1}, β₂={model1.beta2})")
print(f"Total parameters: {model1.alpha.size + model1.alpha0.size + model1.beta.size + 1}")
print()

```

```

print("Training Model 1...")
start_time1 = time.time()
history1 = model1.train(
    X_train, y_train, X_val, y_val,
    epochs=5000,
    print_every=1000,
    verbose=True
)
time1 = time.time() - start_time1

# Evaluate Model 1
y_train_pred1 = model1.predict(X_train)
y_val_pred1 = model1.predict(X_val)
train_acc1 = np.mean(y_train_pred1 == y_train)
val_acc1 = np.mean(y_val_pred1 == y_val)
final_train_loss1 = history1['train_loss'][max(history1['train_loss'].keys())]
final_val_loss1 = history1['val_loss'][max(history1['val_loss'].keys())]

print()
print("Model 1 Results:")
print(f" Training Time: {time1:.2f}s")
print(f" Training Loss: {final_train_loss1:.5f} | Accuracy: {train_acc1:.4f}")
print(f" Validation Loss: {final_val_loss1:.5f} | Accuracy: {val_acc1:.4f}")
print()

# =====
# Model 2: Adam Fully Optimized (Mini-Batch)
# =====
print("=" * 80)
print("MODEL 2: Adam Optimizer (Mini-Batch + Optimizations)")
print("=" * 80)

model2 = NeuralNetworkAdamFullyOptimized(
    input_dim=2,
    hidden_dim=128,

```

```

        learning_rate=0.001,
        beta1=0.9,
        beta2=0.999,
        batch_size=256
    )

    print(f"Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid")
    print(f"Optimizer: Adam (lr={model2.lr}, β₁={model2.beta1}, β₂={model2.beta2})")
    print(f"Batch Size: {model2.batch_size}")
    print(f"Features: Mini-batch, Early Stopping, LR Scheduling, Gradient Clipping")
    print(f"Total parameters: {model2.alpha.size + model2.alpha0.size + model2.beta.size + 1}")
    print()

    print("Training Model 2...")
    start_time2 = time.time()
    history2 = model2.train(
        X_train, y_train, X_val, y_val,
        epochs=5000,
        print_every=1000,
        early_stopping_patience=50,
        lr_schedule='cosine',
        gradient_clip=1.0,
        verbose=True
    )
    time2 = time.time() - start_time2

    # Evaluate Model 2
    y_train_pred2 = model2.predict(X_train)
    y_val_pred2 = model2.predict(X_val)
    train_acc2 = np.mean(y_train_pred2 == y_train)
    val_acc2 = np.mean(y_val_pred2 == y_val)
    final_train_loss2 = history2['train_loss'][max(history2['train_loss'].keys())]

```

```

final_val_loss2 = history2['val_loss'][max(history2['val_loss'].keys())]

print()
print("Model 2 Results:")
print(f" Training Time: {time2:.2f}s")
print(f" Training Loss: {final_train_loss2:.5f} | Accuracy: {train_acc2:.4f}")
print(f" Validation Loss: {final_val_loss2:.5f} | Accuracy: {val_acc2:.4f}")
print()

# =====
# Model 3: SGD with Momentum (Mini-Batch)
# =====
print("=" * 80)
print("MODEL 3: SGD with Momentum (Mini-Batch)")
print("=" * 80)

model3 = NeuralNetworkSGDMomentum(
    input_dim=2,
    hidden_dim=128,
    learning_rate=0.01,
    momentum=0.9,
    batch_size=256
)

print(f"Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid")
print(f"Optimizer: SGD with Momentum (lr={model3.lr}, γ={model3.momentum})")
print(f"Batch Size: {model3.batch_size}")
print(f"Total parameters: {model3.alpha.size + model3.alpha0.size + model3.beta.size + 1}")
print()

print("Training Model 3...")
start_time3 = time.time()
history3 = model3.train()

```

```

        X_train, y_train, X_val, y_val,
        epochs=5000,
        print_every=1000,
        verbose=True
    )
time3 = time.time() - start_time3

# Evaluate Model 3
y_train_pred3 = model3.predict(X_train)
y_val_pred3 = model3.predict(X_val)
train_acc3 = np.mean(y_train_pred3 == y_train)
val_acc3 = np.mean(y_val_pred3 == y_val)
final_train_loss3 = history3['train_loss'][max(history3['train_loss'].keys())]
final_val_loss3 = history3['val_loss'][max(history3['val_loss'].keys())]

print()
print("Model 3 Results:")
print(f" Training Time: {time3:.2f}s")
print(f" Training Loss: {final_train_loss3:.5f} | Accuracy: {train_acc3:.4f}")
print(f" Validation Loss: {final_val_loss3:.5f} | Accuracy: {val_acc3:.4f}")
print()

# =====
# Performance Comparison Summary
# =====
print("=" * 80)
print("PERFORMANCE COMPARISON SUMMARY")
print("=" * 80)
print()
print(f"{'Metric':<30} {'Model 1 (Adam FB)':<20} {'Model 2 (Adam MB)':<20} {'Model 3 (SGD+Mom)':<20}")
print("-" * 90)
print(f"{'Training Time':<30} {time1:>15.2f}s    {time2:>15.2f}s    {time3:>15.2f}s")
print(f"{'Epochs Completed':<30} {max(history1['train_loss'].keys())+1:>20} {max(history2['train_loss'].keys())+1:>20} {max(history3['train_loss'].keys())+1:>20}")

```

```

1:>20}")  

    print(f"{'Final Train Loss':<30} {final_train_loss1:>20.5f} {final_train_loss2:>  

20.5f} {final_train_loss3:>20.5f}")  

    print(f"{'Final Val Loss':<30} {final_val_loss1:>20.5f} {final_val_loss2:>20.5f}  

{final_val_loss3:>20.5f}")  

    print(f"{'Train Accuracy':<30} {train_acc1:>20.4f} {train_acc2:>20.4f} {train  

_acc3:>20.4f}")  

    print(f"{'Val Accuracy':<30} {val_acc1:>20.4f} {val_acc2:>20.4f} {val_acc3:  

>20.4f}")  

    print()  

# =====  

# Visualization  

# =====  

print("Creating comprehensive visualizations...")  

fig = plt.figure(figsize=(20, 12))  

gs = fig.add_gridspec(3, 3, hspace=0.35, wspace=0.35)  

# Row 1: Loss curves  

ax1 = fig.add_subplot(gs[0, :2])  

epochs1 = list(history1['train_loss'].keys())  

epochs2 = list(history2['train_loss'].keys())  

epochs3 = list(history3['train_loss'].keys())  

# Training losses  

ax1.plot(epochs1, list(history1['train_loss'].values()),  

         label='Model 1 Train (Adam FB)', linewidth=2, linestyle='-', alpha=0.8,  

         color='#1f77b4')  

ax1.plot(epochs2, list(history2['train_loss'].values()),  

         label='Model 2 Train (Adam MB)', linewidth=2, linestyle='-', alpha=0.8,  

         color='ff7f0e')  

ax1.plot(epochs3, list(history3['train_loss'].values()),  

         label='Model 3 Train (SGD+Mom)', linewidth=2, linestyle='-', alpha=0.  

8, color='#2ca02c')

```

```

# Validation losses
ax1.plot(epochs1, list(history1['val_loss'].values()),
         label='Model 1 Val', linewidth=2, linestyle='--', alpha=0.8, color='#1f77
b4')
ax1.plot(epochs2, list(history2['val_loss'].values()),
         label='Model 2 Val', linewidth=2, linestyle='--', alpha=0.8, color='#ff7f
0e')
ax1.plot(epochs3, list(history3['val_loss'].values()),
         label='Model 3 Val', linewidth=2, linestyle='--', alpha=0.8, color='#2ca
02c')

ax1.set_xlabel('Epoch', fontsize=12)
ax1.set_ylabel('Binary Cross-Entropy Loss', fontsize=12)
ax1.set_title('Training Progress: All Models Comparison', fontsize=14, fontwe
ight='bold')
ax1.legend(fontsize=10, loc='upper right', ncol=2)
ax1.grid(True, alpha=0.3)
ax1.set_ylim(bottom=0)

# Performance metrics table
ax_table = fig.add_subplot(gs[0, 2])
ax_table.axis('off')

table_data = [
    ['Metric', 'M1', 'M2', 'M3'],
    ['Time (s)', f'{time1:.1f}', f'{time2:.1f}', f'{time3:.1f}'],
    ['Epochs', f'{max(epochs1)+1}', f'{max(epochs2)+1}', f'{max(epochs3)+
1}'],
    ['Train Acc', f'{train_acc1:.4f}', f'{train_acc2:.4f}', f'{train_acc3:.4f}'],
    ['Val Acc', f'{val_acc1:.4f}', f'{val_acc2:.4f}', f'{val_acc3:.4f}'],
    ['Train Loss', f'{final_train_loss1:.4f}', f'{final_train_loss2:.4f}', f'{final_train
_loss3:.4f}'],
    ['Val Loss', f'{final_val_loss1:.4f}', f'{final_val_loss2:.4f}', f'{final_val_loss3:.
4f}']
]

```

```

table = ax_table.table(cellText=table_data, cellLoc='center', loc='center',
                      colWidths=[0.35, 0.22, 0.22, 0.22])
table.auto_set_font_size(False)
table.set_fontsize(9)
table.scale(1, 2)

# Header styling
for i in range(4):
    table[(0, i)].set_facecolor('#4CAF50')
    table[(0, i)].set_text_props(weight='bold', color='white')

# Alternate row colors
for i in range(1, len(table_data)):
    for j in range(4):
        if i % 2 == 0:
            table[(i, j)].set_facecolor('#f0f0f0')

ax_table.set_title('Performance Metrics', fontsize=12, fontweight='bold', pad
=20)

# Row 2: Decision boundaries
ax2 = fig.add_subplot(gs[1, 0])
plot_decision_boundary(ax2, model1, X_val, y_val, 'Model 1: Adam (Full-Batc
h)', val_acc1)

ax3 = fig.add_subplot(gs[1, 1])
plot_decision_boundary(ax3, model2, X_val, y_val, 'Model 2: Adam (Mini-Bat
ch)', val_acc2)

ax4 = fig.add_subplot(gs[1, 2])
plot_decision_boundary(ax4, model3, X_val, y_val, 'Model 3: SGD + Moment
um', val_acc3)

# Row 3: Validation loss comparison (zoomed)
ax5 = fig.add_subplot(gs[2, 0])
ax5.plot(epochs1, list(history1['val_loss'].values())),

```

```

        label='Adam (Full-Batch)', linewidth=2, color='#1f77b4')
ax5.plot(epochs2, list(history2['val_loss'].values())),
        label='Adam (Mini-Batch)', linewidth=2, color='#ff7f0e')
ax5.plot(epochs3, list(history3['val_loss'].values())),
        label='SGD + Momentum', linewidth=2, color='#2ca02c')
ax5.set_xlabel('Epoch', fontsize=11)
ax5.set_ylabel('Validation Loss', fontsize=11)
ax5.set_title('Validation Loss Comparison', fontsize=12, fontweight='bold')
ax5.legend(fontsize=10)
ax5.grid(True, alpha=0.3)

# Training loss comparison (zoomed)
ax6 = fig.add_subplot(gs[2, 1])
ax6.plot(epochs1, list(history1['train_loss'].values())),
        label='Adam (Full-Batch)', linewidth=2, color='#1f77b4')
ax6.plot(epochs2, list(history2['train_loss'].values())),
        label='Adam (Mini-Batch)', linewidth=2, color='#ff7f0e')
ax6.plot(epochs3, list(history3['train_loss'].values())),
        label='SGD + Momentum', linewidth=2, color='#2ca02c')
ax6.set_xlabel('Epoch', fontsize=11)
ax6.set_ylabel('Training Loss', fontsize=11)
ax6.set_title('Training Loss Comparison', fontsize=12, fontweight='bold')
ax6.legend(fontsize=10)
ax6.grid(True, alpha=0.3)

# Accuracy comparison
ax7 = fig.add_subplot(gs[2, 2])
models = ['Adam\n(Full-Batch)', 'Adam\n(Mini-Batch)', 'SGD +\nMomentum']
train_accs = [train_acc1, train_acc2, train_acc3]
val_accs = [val_acc1, val_acc2, val_acc3]

x = np.arange(len(models))
width = 0.35

bars1 = ax7.bar(x - width/2, train_accs, width, label='Train Accuracy',
                 color='#1f77b4', alpha=0.8)

```

```

bars2 = ax7.bar(x + width/2, val_accs, width, label='Val Accuracy',
                 color='#ff7f0e', alpha=0.8)

ax7.set_ylabel('Accuracy', fontsize=11)
ax7.set_title('Final Accuracy Comparison', fontsize=12, fontweight='bold')
ax7.set_xticks(x)
ax7.set_xticklabels(models, fontsize=9)
ax7.legend(fontsize=10)
ax7.grid(True, alpha=0.3, axis='y')
ax7.set_ylim([0.8, 1.0])

# Add value labels on bars
for bars in [bars1, bars2]:
    for bar in bars:
        height = bar.get_height()
        ax7.annotate(f'{height:.3f}',
                     xy=(bar.get_x() + bar.get_width() / 2, height),
                     xytext=(0, 3),
                     textcoords="offset points",
                     ha='center', va='bottom', fontsize=8)

plt.suptitle('Neural Network Optimizer Comparison: Adam vs SGD with Momentum',
             fontsize=16, fontweight='bold', y=0.995)

# Save figure
plt.savefig('hw1_results.png', dpi=150, bbox_inches='tight')
print("Visualization saved as 'hw1_results.png'")
print()

print("=" * 80)
print("Analysis Complete!")
print("=" * 80)

```

```
if __name__ == "__main__":
    main()
```

SGD with Momentum

```
import numpy as np
from typing import Dict
```

```
class NeuralNetworkSGDMomentum:
```

```
    """
```

```
        Two-layer neural network with ReLU activation and SGD with Momentum op
        timizer
```

Architecture:

Input(2) → Hidden(d) → ReLU → Output(1) → Sigmoid

Mathematical formulation:

$s_{ik} = \alpha_{k0} + \sum_j \alpha_{kj} \cdot x_{ij}$ (hidden layer)

$h_{ik} = \max(0, s_{ik})$ (ReLU)

$s_i = \beta_0 + \sum_k \beta_k \cdot h_{ik}$ (output layer)

$p_i = \sigma(s_i) = 1/(1+e^{-s_i})$ (sigmoid)

```
    """
```

```
def __init__(self, input_dim: int = 2, hidden_dim: int = 128,
            learning_rate: float = 0.01, momentum: float = 0.9,
            batch_size: int = 256):
```

```
    """
```

Initialize network parameters and SGD with Momentum optimizer state

Args:

input_dim: Number of input features (p)

hidden_dim: Number of hidden units (d)

learning_rate: Learning rate (η)

momentum: Momentum coefficient (γ), typically 0.9

batch_size: Size of mini-batches for training

```

"""
self.p = input_dim
self.d = hidden_dim
self.lr = learning_rate
self.momentum = momentum
self.batch_size = batch_size

# =====
# Initialize Network Parameters
# =====
# Hidden layer ( $\alpha$  parameters)
self.alpha = np.random.randn(hidden_dim, input_dim) * np.sqrt(2.0 / input
_dim)
self.alpha0 = np.zeros(hidden_dim)

# Output layer ( $\beta$  parameters)
self.beta = np.random.randn(hidden_dim) * np.sqrt(2.0 / hidden_dim)
self.beta0 = 0.0

# =====
# Initialize Momentum State
# =====
# Velocity (momentum buffer) - v
self.v_alpha = np.zeros_like(self.alpha)
self.v_alpha0 = np.zeros_like(self.alpha0)
self.v_beta = np.zeros_like(self.beta)
self.v_beta0 = 0.0

# Cache for backpropagation
self.cache = {}

def forward(self, X: np.ndarray) → np.ndarray:
"""
Forward pass through the network

```

Mathematical operations:

```

1. s_hidden =  $\alpha_0 + \alpha @ X^T$       (hidden pre-activation)
2. h = max(0, s_hidden)           (ReLU activation)
3. s_output =  $\beta_0 + \beta^T @ h$     (output pre-activation)
4. prob =  $1 / (1 + \exp(-s_{\text{output}}))$  (sigmoid activation)

```

Args:

X: Input data of shape (n, p)

Returns:

```

prob: Probabilities of shape (n,) where prob[i] = P(y=1 | x_i)
"""
n = X.shape[0]

```

Step 1: Hidden layer pre-activation

```
s_hidden = self.alpha0[:, np.newaxis] + self.alpha @ X.T # (d, n)
```

Step 2: ReLU activation

```
h = np.maximum(0, s_hidden) # (d, n)
```

Step 3: Output layer pre-activation

```
s_output = self.beta0 + self.beta @ h # (n,)
```

Step 4: Sigmoid activation

```
s_output_clipped = np.clip(s_output, -500, 500) # Numerical stability
prob = 1.0 / (1.0 + np.exp(-s_output_clipped)) # (n,)
```

Cache intermediate values for backpropagation

```
self.cache = {
    'X': X,
    's_hidden': s_hidden,
    'h': h,
    's_output': s_output,
    'prob': prob,
    'n': n
}
```

```

    return prob

def loss_fn(self, y_true: np.ndarray, y_pred: np.ndarray = None) → float:
    """
    Binary cross-entropy loss (negative log-likelihood)

```

Mathematical formula:

$$J = -(1/n) \sum_i [y_i \cdot \log(p_i) + (1-y_i) \cdot \log(1-p_i)]$$

Args:

`y_true`: True labels of shape (n,
`y_pred`: Predicted probabilities of shape (n,) [optional, uses cache if No
ne]

Returns:

`loss`: Scalar binary cross-entropy loss

"""

if `y_pred` is None:

Use numerically stable version from logits

`s` = `self.cache['s_output']`

`loss` = `np.mean(`

`np.maximum(0, s) - s * y_true + np.log(1 + np.exp(-np.abs(s)))`

`)`

else:

`eps` = `1e-15`

`y_pred_clipped` = `np.clip(y_pred, eps, 1 - eps)`

`loss` = `-np.mean(`

`y_true * np.log(y_pred_clipped) +`

`(1 - y_true) * np.log(1 - y_pred_clipped)`

`)`

`return loss`

```

def backward(self, y_true: np.ndarray) → Dict[str, np.ndarray]:
    """

```

Backpropagation: Compute gradients using chain rule

Mathematical derivation:

Step 1: Gradient at output (sigmoid + cross-entropy)

$$\frac{\partial J}{\partial s_{\text{output}}} = p - y$$

Step 2: Gradients for β parameters

$$\frac{\partial J}{\partial \beta_k} = (1/n) \sum_i (p_i - y_i) \cdot h_{ik}$$

$$\frac{\partial J}{\partial \beta_0} = (1/n) \sum_i (p_i - y_i)$$

Step 3: Backprop through output layer

$$\frac{\partial J}{\partial h_{ik}} = (p_i - y_i) \cdot \beta_k$$

Step 4: Gradient through ReLU

$$\frac{\partial J}{\partial s_{ik}} = \frac{\partial J}{\partial h_{ik}} \cdot \mathbb{1}(s_{ik} > 0)$$

Step 5: Gradients for α parameters

$$\frac{\partial J}{\partial \alpha_{kj}} = (1/n) \sum_i (\frac{\partial J}{\partial s_{ik}}) \cdot x_{ij}$$

$$\frac{\partial J}{\partial \alpha_{k0}} = (1/n) \sum_i (\frac{\partial J}{\partial s_{ik}})$$

Args:

`y_true`: True labels of shape (n,)

Returns:

`gradients`: Dictionary containing all parameter gradients

"""

Retrieve cached values from forward pass

X = self.cache['X']

s_hidden = self.cache['s_hidden']

h = self.cache['h']

prob = self.cache['prob']

n = self.cache['n']

Step 1: Gradient at output layer

dJ_ds_output = prob - y_true # (n,)

Step 2: Gradients for β parameters

```

grad_beta = (h @ dJ_ds_output) / n # (d,)
grad_beta0 = np.mean(dJ_ds_output) # scalar

# Step 3: Backprop through output layer
dJ_dh = np.outer(self.beta, dJ_ds_output) # (d, n)

# Step 4: Gradient through ReLU
dh_ds_hidden = (s_hidden > 0).astype(float) # (d, n)
dJ_ds_hidden = dJ_dh * dh_ds_hidden # (d, n)

# Step 5: Gradients for α parameters
grad_alpha = (dJ_ds_hidden @ X) / n # (d, p)
grad_alpha0 = np.mean(dJ_ds_hidden, axis=1) # (d,)

# Return all gradients as dictionary
gradients = {
    'alpha': grad_alpha,
    'alpha0': grad_alpha0,
    'beta': grad_beta,
    'beta0': grad_beta0
}

return gradients

```

def update_parameters(self, gradients: Dict[str, np.ndarray]) → None:

"""

Update parameters using SGD with Momentum optimizer

Momentum Algorithm:

For each parameter θ and its gradient g :

1. $v = \gamma \cdot v + \eta \cdot g$ (update velocity with momentum)
2. $\theta = \theta - v$ (parameter update)

Intuition:

- v (velocity) accumulates gradient history

- γ (momentum) determines how much past gradients influence current update

- Helps accelerate SGD in relevant directions and dampen oscillations
- Typical values: $\gamma \in [0.5, 0.9, 0.99]$

Mathematical Note:

Momentum can be viewed as exponentially weighted moving average:

$$\begin{aligned}v_t &= \gamma \cdot v_{t-1} + \eta \cdot g_t \\&= \eta \cdot \sum_i \gamma^i \cdot g_{t-i}\end{aligned}$$

Args:

gradients: Dictionary of gradients from backward pass

"""

Extract gradients

```
grad_alpha = gradients['alpha']
grad_alpha0 = gradients['alpha0']
grad_beta = gradients['beta']
grad_beta0 = gradients['beta0']
```

=====

Update α parameters (hidden layer weights)

=====

Update velocity: $v = \gamma \cdot v + \eta \cdot g$

```
self.v_alpha = self.momentum * self.v_alpha + self.lr * grad_alpha
```

Update parameter: $\theta = \theta - v$

```
self.alpha -= self.v_alpha
```

=====

Update α_0 parameters (hidden layer biases)

=====

```
self.v_alpha0 = self.momentum * self.v_alpha0 + self.lr * grad_alpha0
```

```
self.alpha0 -= self.v_alpha0
```

=====

Update β parameters (output layer weights)

```

# =====
self.v_beta = self.momentum * self.v_beta + self.lr * grad_beta
self.beta -= self.v_beta

# =====
# Update  $\beta_0$  parameter (output layer bias)
# =====
self.v_beta0 = self.momentum * self.v_beta0 + self.lr * grad_beta0
self.beta0 -= self.v_beta0

def train(self, X_train: np.ndarray, y_train: np.ndarray,
          X_val: np.ndarray = None, y_val: np.ndarray = None,
          epochs: int = 50000, print_every: int = 1000,
          verbose: bool = True) → Dict[str, Dict[int, float]]:
"""
Training loop with mini-batch SGD and momentum

```

Algorithm:

For each epoch:

1. Shuffle training data
2. For each mini-batch:
 - a. Forward pass: compute predictions
 - b. Compute loss: measure error
 - c. Backward pass: compute gradients
 - d. Update parameters: apply SGD with momentum
3. (Optional) Validate and log progress

Args:

- X_train: Training inputs of shape (n_train, p)
- y_train: Training labels of shape (n_train,)
- X_val: Validation inputs of shape (n_val, p) [optional]
- y_val: Validation labels of shape (n_val,) [optional]
- epochs: Number of training iterations
- print_every: Frequency of progress logging
- verbose: Whether to print progress

Returns:

```
    history: Dictionary containing training and validation losses
"""
n_samples = X_train.shape[0]
n_batches = int(np.ceil(n_samples / self.batch_size))

history = {
    'train_loss': {},
    'val_loss': {}
}

for epoch in range(epochs):
    # Shuffle data for stochastic gradient descent
    indices = np.random.permutation(n_samples)
    X_shuffled = X_train[indices]
    y_shuffled = y_train[indices]

    epoch_losses = []

    # =====
    # Mini-batch Training
    # =====
    for batch_idx in range(n_batches):
        start_idx = batch_idx * self.batch_size
        end_idx = min(start_idx + self.batch_size, n_samples)

        X_batch = X_shuffled[start_idx:end_idx]
        y_batch = y_shuffled[start_idx:end_idx]

        # 1. Forward pass
        y_pred = self.forward(X_batch)

        # 2. Compute loss
        batch_loss = self.loss_fn(y_batch)
        epoch_losses.append(batch_loss)
```

```

# 3. Backward pass
gradients = self.backward(y_batch)

# 4. Update parameters
self.update_parameters(gradients)

# Average loss for the epoch
train_loss = np.mean(epoch_losses)
history['train_loss'][epoch] = train_loss

# =====
# Validation Step (if validation data provided)
# =====
if X_val is not None and y_val is not None:
    y_val_pred = self.forward(X_val)
    val_loss = self.loss_fn(y_val, y_val_pred)
    history['val_loss'][epoch] = val_loss
else:
    val_loss = None

# =====
# Logging
# =====
if verbose and (epoch % print_every == 0 or epoch == epochs - 1):
    log_str = f"Epoch {epoch:05d}/{epochs:05d} | Train Loss: {train_loss:.5f}"
    if val_loss is not None:
        log_str += f" | Val Loss: {val_loss:.5f}"

    # Compute accuracy
    train_acc = np.mean((self.forward(X_train) > 0.5) == y_train)
    val_acc = np.mean((y_val_pred > 0.5) == y_val)
    log_str += f" | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f}"

print(log_str)

```

```
    return history
```

```
def predict(self, X: np.ndarray, threshold: float = 0.5) → np.ndarray:
```

```
    """
```

```
        Make binary predictions
```

Mathematical operation:

1. prob = forward(X) → compute $P(y=1 | x)$
2. prediction = $\mathbb{1}(\text{prob} > \text{threshold})$

Args:

X: Input data of shape (n, p)

threshold: Decision threshold (default 0.5)

Returns:

predictions: Binary predictions of shape (n,)

```
"""
```

```
prob = self.forward(X)
```

```
predictions = (prob > threshold).astype(int)
```

```
return predictions
```

```
def predict_proba(self, X: np.ndarray) → np.ndarray:
```

```
    """
```

```
        Predict probabilities  $P(y=1 | x)$ 
```

Args:

X: Input data of shape (n, p)

Returns:

probabilities: Probabilities of shape (n,)

```
"""
```

```
return self.forward(X)
```

Performance benchmark result

=====

=====

STATS 413 - HW1: Neural Network Optimizer Comparison

Task: Learn quarter circle decision boundary

=====

=====

Generating data...

Training samples: 10000

Validation samples: 5000

Positive samples (inside circle): 7906/10000 = 79.06%

=====

=====

MODEL 1: Adam Optimizer (Full-Batch)

=====

=====

Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid

Optimizer: Adam ($\text{lr}=0.001$, $\beta_1=0.9$, $\beta_2=0.999$)

Total parameters: 513

Training Model 1...

Epoch 0000000/0005000 | Train Loss: 0.75979 | Val Loss: 0.74669 | Train Acc: 0.4166 | Val Acc: 0.4410

Epoch 0001000/0005000 | Train Loss: 0.10228 | Val Loss: 0.10032 | Train Acc: 0.9665 | Val Acc: 0.9674

Epoch 0002000/0005000 | Train Loss: 0.05742 | Val Loss: 0.05674 | Train Acc: 0.9888 | Val Acc: 0.9902

Epoch 0003000/0005000 | Train Loss: 0.03762 | Val Loss: 0.03747 | Train Acc: 0.9945 | Val Acc: 0.9936

Epoch 0004000/0005000 | Train Loss: 0.02726 | Val Loss: 0.02723 | Train Acc: 0.9945 | Val Acc: 0.9940

Epoch 0004999/0005000 | Train Loss: 0.02072 | Val Loss: 0.02082 | Train Acc: 0.9961 | Val Acc: 0.9960

Model 1 Results:

Training Time: 338.40s

Training Loss: 0.02072 | Accuracy: 0.9961

Validation Loss: 0.02082 | Accuracy: 0.9960

=====

=====

MODEL 2: Adam Optimizer (Mini-Batch + Optimizations)

=====

=====

Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid

Optimizer: Adam ($\text{lr}=0.001$, $\beta_1=0.9$, $\beta_2=0.999$)

Batch Size: 256

Features: Mini-batch, Early Stopping, LR Scheduling, Gradient Clipping

Total parameters: 513

Training Model 2...

Epoch 00000 | Train: 0.62285 | Val: 0.57163 | LR: 0.001000

Epoch 01000 | Train: 0.01586 | Val: 0.01601 | LR: 0.000905

Early stopping at epoch 1651

Best val loss: 0.01302

Model 2 Results:

Training Time: 78.13s

Training Loss: 0.01242 | Accuracy: 0.9972

Validation Loss: 0.01413 | Accuracy: 0.9970

=====

=====

MODEL 3: SGD with Momentum (Mini-Batch)

=====

=====

Architecture: Input(2) → Hidden(128) → ReLU → Output(1) → Sigmoid

Optimizer: SGD with Momentum ($\text{lr}=0.01$, $\gamma=0.9$)

Batch Size: 256

Total parameters: 513

Training Model 3...

Epoch 00000/05000 | Train Loss: 0.60257 | Val Loss: 0.52839 | Train Acc: 0.7906 | Val Acc: 0.7966

Epoch 01000/05000 | Train Loss: 0.02012 | Val Loss: 0.02002 | Train Acc: 0.9962 | Val Acc: 0.9964

Epoch 02000/05000 | Train Loss: 0.01409 | Val Loss: 0.01437 | Train Acc: 0.9974 | Val Acc: 0.9972

Epoch 03000/05000 | Train Loss: 0.01244 | Val Loss: 0.01178 | Train Acc: 0.9974 | Val Acc: 0.9976

Epoch 04000/05000 | Train Loss: 0.01041 | Val Loss: 0.01104 | Train Acc: 0.9978 | Val Acc: 0.9964

Epoch 04999/05000 | Train Loss: 0.00880 | Val Loss: 0.00947 | Train Acc: 0.9989 | Val Acc: 0.9984

Model 3 Results:

Training Time: 161.37s

Training Loss: 0.00880 | Accuracy: 0.9989

Validation Loss: 0.00947 | Accuracy: 0.9984

=====

=====

PERFORMANCE COMPARISON SUMMARY

=====

=====

Metric	Model 1 (Adam FB)	Model 2 (Adam MB)	Model 3 (SG D+Mom)
--------	-------------------	-------------------	--------------------

Training Time	338.40s	78.13s	161.37s
---------------	---------	--------	---------

Epochs Completed	5000	1652	5000
------------------	------	------	------

Final Train Loss	0.02072	0.01242	0.00880
------------------	---------	---------	---------

Final Val Loss	0.02082	0.01413	0.00947
----------------	---------	---------	---------

Train Accuracy	0.9961	0.9972	0.9989
----------------	--------	--------	--------

Val Accuracy	0.9960	0.9970	0.9984
--------------	--------	--------	--------

Creating comprehensive visualizations...
Visualization saved as 'hw1_results.png'

=====

=====

Analysis Complete!

=====

=====