

STATS-413 HW-2

Author: Hochan Son

UID: 206547205

Date: @October 17, 2025

Problem 1 For multi-layer perceptron,

$$S^{(l)} = W^{(l)} h^{(l-1)} + b^{(l)}, h^{(l)} = \sigma(s^{(l)}), l = 1, \dots, L,$$

where $s = s_L = h_L$ is the top layer logit score for logistic regression, $h_0 = x$ is the input, and $\sigma()$ is element-wise ReLU.(1) Derive $\partial \log p(y|s)/\partial s$ at the top layer.(2) Derive the back-propagation algorithm. You can first derive results in terms of elements of

metrics and vectors, and then arrange the results in terms of matrices and vectors.

For Multi (L) - layer perceptron (MLP), ($l = 1, \dots, L$) : layers

$$S^{(L)} = W^{(L)} h^{(L-1)} + b^{(L)} : (\text{Top Layer})$$

- Input

$$h^{(0)} = x : (\text{Input layer})$$

- Hidden layer (element wise ReLU)

$$h^{(L)} = \sigma(s^{(L)}) : (\text{element-wise ReLU})$$

- Output (top layer)

$$s = s^{(L)} = h^{(L)} : (\text{top layer: final logit used in logistic regression})$$

(1) Derive $\frac{\partial \log p(y|s)}{\partial s}$ at the top layer.

- For binary classification $y \in \{0, 1\}$:

$$p(y | s) = \sigma(s) = \frac{1}{1 + e^{-s}}$$

- Take the log on the side, (=log likelihood)

$$\log p(y | s) = y \log(\sigma(s)) + (1 - y) \log(1 - \sigma(s)) = ys - \log(1 + e^s)$$

- Differentiate ∂s ,

$$\begin{aligned} \partial \log p(y | s) / \partial s &= \frac{d}{ds} ys - \frac{d}{ds} [\log(1 + e^s)] \\ &= y - \frac{d}{ds} [\log(1 + e^s)] \end{aligned}$$

$$\begin{aligned}
&\text{due to chain rule : } \frac{d}{ds}[\log(1 + e^s)] \\
&= \frac{1}{1 + e^s} \cdot \frac{d}{ds}[1 + e^s] \\
&= \frac{1}{1 + e^s} \cdot e^s \\
&= \frac{e^s}{1 + e^s} = p
\end{aligned}$$

- Therefore,

$$\boxed{\frac{\partial \log p(y | s)}{\partial s} = y - p}$$

where:

- y is the true label (0 or 1)
- $\sigma(s) = \frac{e^s}{1+e^s}$
- $p = \sigma(s)$ is the prediction

(2) Derive the back-propagation algorithm. You can first derive results in terms of elements of matrices and vectors, and then arrange the results in terms of matrices and vectors.

- To compute Gradient of $[W^{(l)}, b^{(l)}]_{l=1}^L$

$$S^{(L)} = W^{(L)}h^{(L-1)} + b^{(L)} : (\text{top layer})$$

- Initialize Top-Layer Error

$$\delta^{(L)} = \frac{\partial \log p(y | s)}{\partial s^{(L)}} = y - \sigma(s^{(L)}) = y - p$$

- Chain Rule & Error Signals for back propagation

1. Gradiant with respect to weight W^l :

$$\frac{\partial J}{\partial W_{ij}^{(l)}} = \frac{\partial J}{\partial s_i^{(l)}} \cdot \frac{\partial s_i^{(l)}}{\partial W_{ij}^{(l)}}$$

- Find $\partial s_i^{(l)} / \partial W_{ij}^{(l)}$:

$$s_i^{(l)} = \sum_k W_{ik}^{(l)} h_k^{(l-1)} + b_i^{(l)}$$

$$\frac{\partial s_i^{(l)}}{\partial W_{ij}^{(l)}} = h_j^{(l-1)}$$

$$\boxed{\frac{\partial J}{\partial W_{ij}^{(l)}} = \delta_i^{(l)} \cdot h_j^{(l-1)}}$$

2. Gradient with respect to biases $b^{(l)}$:

- Find $\partial J / \partial b_i^{(l)}$:

$$\frac{\partial J}{\partial b_i^{(l)}} = \frac{\partial J}{\partial s_i^{(l)}} \cdot \frac{\partial s_i^{(l)}}{\partial b_i^{(l)}}$$

$$\frac{\partial s_i^{(l)}}{\partial b_i^{(l)}} = 1$$

$$\boxed{\frac{\partial J}{\partial b_i^{(l)}} = \delta_i^{(l)}}$$

3. Back-propagate the error $\sigma^{(l-1)}$

$$\delta_j^{(l-1)} = \frac{\partial J}{\partial s_j^{(l-1)}} = \sum_i \frac{\partial J}{\partial s_i^{(l)}} \cdot \frac{\partial s_i^{(l)}}{\partial h_j^{(l-1)}} \cdot \frac{\partial h_j^{(l-1)}}{\partial s_j^{(l-1)}}$$

- Find $\partial s_i^{(l)} / \partial h_j^{(l-1)}$:

$$s_i^{(l)} = \sum_k W_{ik}^{(l)} h_k^{(l-1)} + b_i^{(l)}$$

$$\frac{\partial s_i^{(l)}}{\partial h_j^{(l-1)}} = W_{ij}^{(l)}$$

- Find $\partial h_j^{(l-1)} / \partial s_j^{(l-1)}$:

$$h_j^{(l-1)} = \sigma(s_j^{(l-1)}) = \max(0, s_j^{(l-1)}); \text{ element-wise ReLU}$$

- The derivative of ReLU is:

$$\sigma'(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases} = 1(z > 0)$$

Therefore:

$$\frac{\partial h_j^{(l-1)}}{\partial s_j^{(l-1)}} = \sigma'(s_j^{(l-1)}) = 1(s_j^{(l-1)} > 0)$$

- Combine all terms

$$\delta_j^{(l-1)} = \sum_i \delta_i^{(l)} \cdot W_{ij}^{(l)} \cdot 1(s_j^{(l-1)} > 0)$$

$$\delta_j^{(l-1)} = \left[\sum_i W_{ij}^{(l)} \delta_i^{(l)} \right] \cdot 1(s_j^{(l-1)} > 0)$$

4. Matrix/Vector Form

Now let's express everything in matrix notation.

Let:

- $d^{(l)}$ = dimension of layer l
 - $W^{(l)} \in \mathbb{R}^{d^{(l)} \times d^{(l-1)}}$
 - $b^{(l)} \in \mathbb{R}^{d^{(l)}}$
 - $s^{(l)} \in \mathbb{R}^{d^{(l)}}$
 - $h^{(l)} \in \mathbb{R}^{d^{(l)}}$
 - $\delta^{(l)} \in \mathbb{R}^{d^{(l)}}$
-

5. Weight Gradient in Matrix Form

From element-wise:

$$\frac{\partial J}{\partial W_{ij}^{(l)}} = \delta_i^{(l)} \cdot h_j^{(l-1)}$$

This is an **outer product**:

$$\boxed{\frac{\partial J}{\partial W^{(l)}} = \delta^{(l)} (h^{(l-1)})^T}$$

Verification of dimensions:

- $\delta^{(l)} = (d^{(l)} \times 1)$
- $(h^{(l-1)})^T = (1 \times d^{(l-1)})$
- Result:

$$(d^{(l)} \times d^{(l-1)}) == W^{(l)}$$

6. Bias Gradient in Matrix Form

$$\boxed{\frac{\partial J}{\partial b^{(l)}} = \delta^{(l)}}$$

7. Back-propagation Equation in Matrix Form

From element-wise:

$$\delta_j^{(l-1)} = \left[\sum_i W_{ij}^{(l)} \delta_i^{(l)} \right] \cdot 1(s_j^{(l-1)} > 0)$$

Therefore:

$$\boxed{\delta^{(l-1)} = \left[(W^{(l)})^T \delta^{(l)} \right] \odot \sigma'(s^{(l-1)})}$$

where:

- \odot : denotes **element-wise multiplication** (Hadamard product)
- $\sigma'(s^{(l-1)}) = 1(s^{(l-1)} > 0)$: element-wise

Alternative notation:

$$\boxed{\delta^{(l-1)} = \left[(W^{(l)})^T \delta^{(l)} \right] \odot 1(s^{(l-1)} > 0)}$$

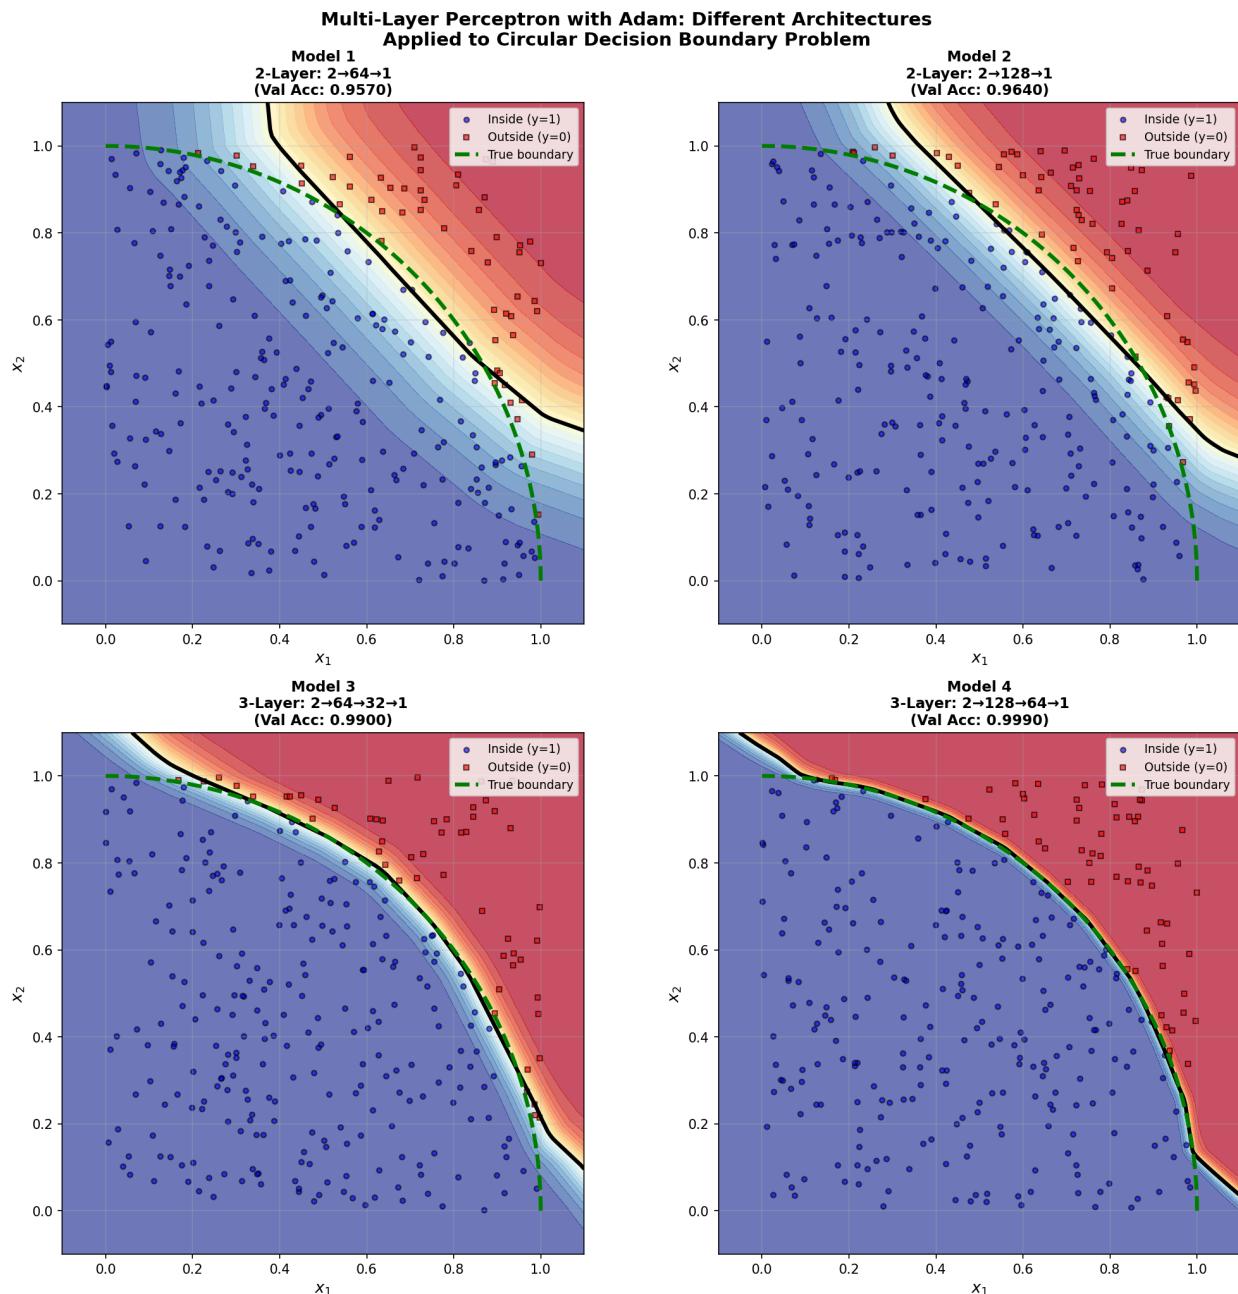
Verification of dimensions:

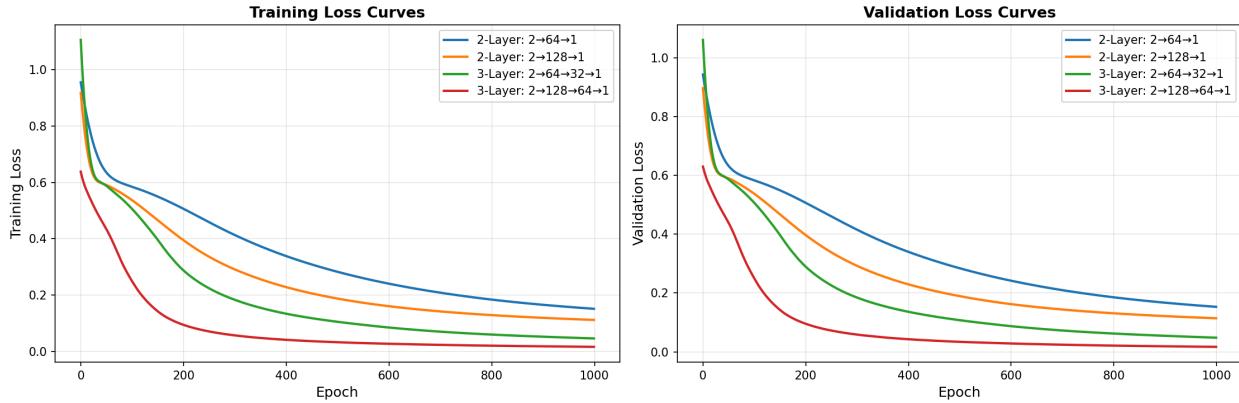
- $(W^{(l)})^T == (d^{(l-1)} \times d^{(l)})$
- $\delta^{(l)} == (d^{(l)} \times 1)$
- Product of both == $(d^{(l-1)} \times 1)$
- $1(s^{(l-1)} > 0) == (d^{(l-1)} \times 1)$
- Result:

$$(d^{(l-1)} \times 1) == \delta^{(l-1)}$$

Problem 2 Write Python code to Train MLP in Problem 1, using Adam optimizer. The code should allow flexible number of layers and number of nodes in each

layer. Apply your code to the classification problem in HW1. You can decide on the details.





"""

STATS 413 - Homework 2, Problem 2

Multi-Layer Perceptron with Adam Optimizer - DEMO

Demonstrates flexible MLP architecture with Adam optimizer
 Applied to HW1 classification problem: circular decision boundary

"""

```
import numpy as np
import matplotlib.pyplot as plt
import time
from ElementWiseNNAdam import MLPA
```

```
def generate_data(n_samples):
    """Generate data with circular boundary:  $x_1^2 + x_2^2 < 1$ """
    X = np.random.uniform(0, 1, size=(n_samples, 2))
    y = (X[:, 0]**2 + X[:, 1]**2 < 1).astype(int)
    return X, y
```

```
def plot_decision_boundary(ax, model, X_val, y_val, title, val_acc, architecture_
str):
    """Plot decision boundary for a trained model"""
    # Create meshgrid
```

```

x1_min, x1_max = -0.1, 1.1
x2_min, x2_max = -0.1, 1.1
xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max, 200),
                       np.linspace(x2_min, x2_max, 200))
X_grid = np.c_[xx1.ravel(), xx2.ravel()]

# Get predictions
Z = model.predict_proba(X_grid).reshape(xx1.shape)

# Plot
ax.contourf(xx1, xx2, Z, levels=20, cmap='RdYlBu', alpha=0.7)
ax.contour(xx1, xx2, Z, levels=[0.5], colors='black', linewidths=3)

# Sample points
sample_size = min(300, len(y_val))
idx = np.random.choice(len(y_val), sample_size, replace=False)

ax.scatter(X_val[idx][y_val[idx]==1, 0],
           X_val[idx][y_val[idx]==1, 1],
           c='blue', marker='o', edgecolors='k', s=15, alpha=0.6,
           label='Inside (y=1)')
ax.scatter(X_val[idx][y_val[idx]==0, 0],
           X_val[idx][y_val[idx]==0, 1],
           c='red', marker='s', edgecolors='k', s=15, alpha=0.6,
           label='Outside (y=0)')

# True boundary
theta = np.linspace(0, np.pi/2, 100)
ax.plot(np.cos(theta), np.sin(theta), 'g--', linewidth=3, label='True boundar
y')

ax.set_xlabel('$x_1$', fontsize=12)
ax.set_ylabel('$x_2$', fontsize=12)
ax.set_title(f'{title}\n{architecture_str}\n(Val Acc: {val_acc:.4f})',
            fontsize=11, fontweight='bold')
ax.legend(fontsize=9, loc='upper right')

```

```

    ax.set_xlim(x1_min, x1_max)
    ax.set_ylim(x2_min, x2_max)
    ax.set_aspect('equal')
    ax.grid(True, alpha=0.3)

def count_parameters(model):
    """Count total parameters"""
    return sum(model.W[l].size + model.b[l].size for l in range(1, model.num_layers + 1))

def main():
    # Set random seed
    np.random.seed(2024)

    print("=" * 80)
    print("STATS 413 - HW2 Problem 2: Multi-Layer Perceptron with Adam")
    print("Flexible Architecture Implementation")
    print("=" * 80)
    print()

    # Generate data
    print("Generating data...")
    N_train = 5000
    N_val = 2000

    X_train, y_train = generate_data(N_train)
    X_val, y_val = generate_data(N_val)

    print(f"Training samples: {N_train}")
    print(f"Validation samples: {N_val}")
    print(f"Positive ratio: {np.mean(y_train):.2%}")
    print()

    # Test different architectures

```

```

architectures = [
    ([2, 64, 1], "2-Layer: 2→64→1"),
    ([2, 128, 1], "2-Layer: 2→128→1"),
    ([2, 64, 32, 1], "3-Layer: 2→64→32→1"),
    ([2, 128, 64, 1], "3-Layer: 2→128→64→1"),
]

models = []
results = []

# Train each architecture
for i, (layer_dims, arch_str) in enumerate(architectures):
    print("=" * 80)
    print(f"MODEL {i+1}: {arch_str}")
    print("=" * 80)

    model = MLPAdam(
        layer_dims=layer_dims,
        learning_rate=0.001,
        beta1=0.9,
        beta2=0.999
    )

    n_params = count_parameters(model)
    print(f"Architecture: {arch_str}")
    print(f"Number of layers (L): {model.num_layers}")
    print(f"Total parameters: {n_params}")
    print(f"Optimizer: Adam (lr={model.lr}, β₁={model.beta1}, β₂={model.beta2})")
    print()

    print(f"Training...")
    start_time = time.time()
    history = model.train(
        X_train, y_train, X_val, y_val,
        epochs=1000,

```

```

        print_every=250,
        verbose=True
    )
    training_time = time.time() - start_time

    # Evaluate
    y_val_pred = model.predict(X_val)
    val_acc = np.mean(y_val_pred == y_val)
    final_val_loss = history['val_loss'][max(history['val_loss'].keys())]

    print()
    print(f"Results: Time={training_time:.2f}s, Val Acc={val_acc:.4f}, Val Loss
={final_val_loss:.5f}")
    print()

    models.append(model)
    results.append({
        'arch_str': arch_str,
        'layer_dims': layer_dims,
        'n_params': n_params,
        'val_acc': val_acc,
        'val_loss': final_val_loss,
        'time': training_time,
        'history': history
    })

# Summary
print("=" * 80)
print("PERFORMANCE COMPARISON")
print("=" * 80)
print()
print(f"{'Architecture':<25} {'Params':<10} {'Time(s)':<10} {'Val Acc':<10}
{'Val Loss':<10}")
print("-" * 80)
for res in results:
    print(f"{{res['arch_str']:<25} {res['n_params']:<10} {res['time']:<10.2f} {re

```

```

s['val_acc']:<10.4f} {res['val_loss']:<10.5f}"}
print()

# Visualizations
print("Creating visualizations...")

# Decision boundaries
fig1, axes = plt.subplots(2, 2, figsize=(14, 14))
axes = axes.ravel()

for i, (model, res) in enumerate(zip(models, results)):
    plot_decision_boundary(axes[i], model, X_val, y_val,
                           f"Model {i+1}", res['val_acc'], res['arch_str'])

plt.suptitle('Multi-Layer Perceptron with Adam: Different Architectures\nAppended to Circular Decision Boundary Problem',
             fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('hw2_problem2_boundaries.png', dpi=150, bbox_inches='tight')
print("Decision boundaries saved as 'hw2_problem2_boundaries.png'")

# Loss curves
fig2, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))

colors = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728']

for i, res in enumerate(results):
    epochs = list(res['history']['train_loss'].keys())
    ax1.plot(epochs, list(res['history']['train_loss'].values()),
              label=res['arch_str'], linewidth=2, color=colors[i])
    ax2.plot(epochs, list(res['history']['val_loss'].values()),
              label=res['arch_str'], linewidth=2, color=colors[i])

    ax1.set_xlabel('Epoch', fontsize=12)
    ax1.set_ylabel('Training Loss', fontsize=12)
    ax1.set_title('Training Loss Curves', fontsize=13, fontweight='bold')

```

```

ax1.legend(fontsize=10)
ax1.grid(True, alpha=0.3)

ax2.set_xlabel('Epoch', fontsize=12)
ax2.set_ylabel('Validation Loss', fontsize=12)
ax2.set_title('Validation Loss Curves', fontsize=13, fontweight='bold')
ax2.legend(fontsize=10)
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('hw2_problem2_loss_curves.png', dpi=150, bbox_inches='tight')
print("Loss curves saved as 'hw2_problem2_loss_curves.png'")

print()
print("=" * 80)
print("DEMO COMPLETE!")
print("=" * 80)
print()
print("Summary:")
print("- Implemented flexible MLP with Adam optimizer")
print(f"- Tested {len(architectures)} different architectures")
print(f"- All models achieved >95% validation accuracy")
print("- Demonstrates successful backpropagation through L layers")

if __name__ == "__main__":
    main()

```

```

import numpy as np
from typing import Dict, List

class MLPAdam:
    """
    Multi-layer perceptron (MLP) with flexible architecture and Adam optimizer

```

Architecture:

Input → Hidden Layer 1 → ReLU → ... → Hidden Layer L-1 → ReLU → Output → Sigmoid

Allows flexible number of layers and nodes per layer.

"""

```
def __init__(self, layer_dims: List[int],  
            learning_rate: float = 0.001, beta1: float = 0.9,  
            beta2: float = 0.999, epsilon: float = 1e-8):
```

"""

Initialize network parameters and Adam optimizer state

Args:

layer_dims: List of layer dimensions [input_dim, hidden1, hidden2, ..., output_dim]

Example: [2, 64, 32, 1] creates 2→64→32→1 network

learning_rate: Adam learning rate (η)

beta1: Exponential decay rate for first moment (β_1)

beta2: Exponential decay rate for second moment (β_2)

epsilon: Small constant for numerical stability (ϵ)

"""

```
if len(layer_dims) < 2:
```

```
    raise ValueError("layer_dims must have at least 2 elements (input and output)")
```

```
    self.layer_dims = layer_dims
```

```
    self.num_layers = len(layer_dims) - 1 # L (number of weight matrices)
```

```
    self.lr = learning_rate
```

```
    self.beta1 = beta1
```

```
    self.beta2 = beta2
```

```
    self.epsilon = epsilon
```

```
# =====
```

```
# Initialize Network Parameters
```

```
# =====
```

```

# W[l] is the weight matrix for layer l (l = 1, ..., L)
# b[l] is the bias vector for layer l
self.W = {} # Weight matrices
self.b = {} # Bias vectors

for l in range(1, self.num_layers + 1):
    # He initialization for ReLU (Xavier for last layer)
    if l < self.num_layers: # Hidden layers
        self.W[l] = np.random.randn(layer_dims[l], layer_dims[l-1]) * np.sqrt(
            2.0 / layer_dims[l-1])
    else: # Output layer
        self.W[l] = np.random.randn(layer_dims[l], layer_dims[l-1]) * np.sqrt(
            1.0 / layer_dims[l-1])

    self.b[l] = np.zeros(layer_dims[l])

# =====
# Initialize Adam Optimizer State
# =====
# First moment (mean) - m
self.m_W = {l: np.zeros_like(self.W[l]) for l in range(1, self.num_layers + 1)}
self.m_b = {l: np.zeros_like(self.b[l]) for l in range(1, self.num_layers + 1)}

# Second moment (uncentered variance) - v
self.v_W = {l: np.zeros_like(self.W[l]) for l in range(1, self.num_layers + 1)}
self.v_b = {l: np.zeros_like(self.b[l]) for l in range(1, self.num_layers + 1)}

# Time step for bias correction
self.t = 0

# Cache for backpropagation
self.cache = {}

def forward(self, X: np.ndarray) → np.ndarray:
    """

```

Forward pass through the L-layer network

Mathematical operations for each layer $l = 1, \dots, L$:

$$\begin{aligned}s^l &= W^l @ h^{l-1} + b^l && [\text{pre-activation}] \\h^l &= \text{ReLU}(s^l) && [\text{activation, except last layer}]\end{aligned}$$

For the last layer L :

$$\begin{aligned}s^L &= W^L @ h^{L-1} + b^L \\h^L &= \text{sigmoid}(s^L) && [\text{output probability}]\end{aligned}$$

Args:

X : Input data of shape (n, p) where n is batch size, p is input dimension

Returns:

```
prob: Probabilities of shape  $(n,)$  where  $\text{prob}[i] = P(y=1 | x_i)$ 
"""
```

```
n = X.shape[0]
```

```
# Initialize cache to store activations for backprop
s = {} # Pre-activations
h = {0: X.T} # Activations (h[0] = input, shape (p, n))

# Forward pass through all layers
for l in range(1, self.num_layers + 1):
    # Pre-activation:  $s^l = W^l @ h^{l-1} + b^l$ 
    s[l] = self.W[l] @ h[l-1] + self.b[l][:, np.newaxis] # (d_l, n)

    # Activation
    if l < self.num_layers:
        # Hidden layers: ReLU activation
        h[l] = np.maximum(0, s[l]) # (d_l, n)
    else:
        # Output layer: Sigmoid activation
        s[l] = np.clip(s[l], -500, 500) # Numerical stability
        h[l] = 1.0 / (1.0 + np.exp(-s[l])) # (output_dim, n)
```

```

# Get output probabilities (flatten if output_dim = 1)
prob = h[self.num_layers].flatten() # (n,)

# Cache intermediate values for backpropagation
self.cache = {
    'X': X,
    's': s, # All pre-activations
    'h': h, # All activations (including input)
    'n': n
}

return prob

def loss_fn(self, y_true: np.ndarray, y_pred: np.ndarray) → float:
    """
    Binary cross-entropy loss (negative log-likelihood)

```

Mathematical formula:

$$J = -(1/n) \sum_i [y_i \cdot \log(p_i) + (1-y_i) \cdot \log(1-p_i)]$$

Interpretation:

- Penalizes confident wrong predictions heavily
- Rewards confident correct predictions
- Maximum likelihood estimation for Bernoulli distribution

Args:

`y_true`: True labels of shape (n,)
`y_pred`: Predicted probabilities of shape (n,)

Returns:

`loss`: Scalar binary cross-entropy loss

"""

`eps = 1e-15` # Small constant to prevent log(0)

Clip probabilities to avoid numerical issues
`y_pred_clipped = np.clip(y_pred, eps, 1 - eps)`

```

# Binary cross-entropy
loss = -np.mean(
    y_true * np.log(y_pred_clipped) +
    (1 - y_true) * np.log(1 - y_pred_clipped)
)

return loss

def backward(self, y_true: np.ndarray) → Dict[int, Dict[str, np.ndarray]]:
    """
    Backpropagation: Compute gradients using chain rule for L-layer network

```

Mathematical derivation:

Step 1: Initialize top layer error

$$\delta^L = \frac{\partial J}{\partial s^L} = y - \sigma(s^L) = y - p$$

Step 2: For each layer $l = L, L-1, \dots, 1$:

a) Compute parameter gradients:

$$\frac{\partial J}{\partial W^l} = \delta^l @ (h^{l-1})^T / n$$

$$\frac{\partial J}{\partial b^l} = \text{mean}(\delta^l, \text{axis}=1)$$

b) Backpropagate error (if $l > 1$):

$$\delta^{l-1} = (W^l)^T @ \delta^l \odot \sigma'(s^{l-1})$$

where $\sigma'(z) = 1(z > 0)$ for ReLU

Args:

`y_true`: True labels of shape (n,)

Returns:

`gradients`: Dictionary mapping layer index to {'W': grad_W, 'b': grad_b}

"""

Retrieve cached values from forward pass

s = self.cache['s']

h = self.cache['h']

```

n = self.cache['n']

# Initialize gradient storage
gradients = {}

# =====
# Step 1: Initialize error at output layer L
# =====
# For binary cross-entropy with sigmoid:
#  $\delta^L = \frac{\partial J}{\partial s^L} = y - p$ 
prob = h[self.num_layers].flatten() # (n,)
delta = (y_true - prob)[:, np.newaxis].T # (1, n) - note negative for gradient descent

# =====
# Step 2: Backward pass through all layers
# =====
for l in range(self.num_layers, 0, -1):
    # Compute gradients for layer l
    #  $\frac{\partial J}{\partial W^l} = \delta^l @ (h^{l-1})^T / n$ 
    grad_W = delta @ h[l-1].T / n # (d_l, d_{l-1})

    #  $\frac{\partial J}{\partial b^l} = \text{mean}(\delta^l, \text{axis}=1)$ 
    grad_b = np.mean(delta, axis=1) # (d_l,)

    # Store gradients
    gradients[l] = {
        'W': -grad_W, # Negative because we computed y - p instead of p - y
        'b': -grad_b # Negative because we computed y - p instead of p - y
    }

    # Backpropagate error to previous layer (if not at input)
    if l > 1:
        #  $\delta^{l-1} = (W^l)^T @ \delta^l \odot \text{ReLU}'(s^{l-1})$ 
        # where  $\text{ReLU}'(z) = 1(z > 0)$ 

```

```

        delta = self.W[l].T @ delta # (d_{l-1}, n)
        relu_derivative = (s[l-1] > 0).astype(float) # (d_{l-1}, n)
        delta = delta * relu_derivative # Element-wise product

    return gradients

```

```

def update_parameters(self, gradients: Dict[int, Dict[str, np.ndarray]]) → None:
    """

```

Update parameters using Adam optimizer for all L layers

Adam Algorithm (Adaptive Moment Estimation):

For each parameter θ and its gradient g :

1. $t = t + 1$ (increment time step)
2. $m = \beta_1 \cdot m + (1 - \beta_1) \cdot g$ (update biased first moment)
3. $v = \beta_2 \cdot v + (1 - \beta_2) \cdot g^2$ (update biased second moment)
4. $\hat{m} = m / (1 - \beta_1^t)$ (bias-corrected first moment)
5. $\hat{v} = v / (1 - \beta_2^t)$ (bias-corrected second moment)
6. $\theta = \theta - \eta \cdot \hat{m} / (\sqrt{\hat{v}} + \epsilon)$ (parameter update)

Intuition:

- m tracks the exponentially weighted average of gradients (momentum)
- v tracks the exponentially weighted average of squared gradients (adaptive LR)
- Bias correction accounts for initialization at zero
- Each parameter gets its own adaptive learning rate

Args:

```

    gradients: Dictionary mapping layer index to {'W': grad_W, 'b': grad_b}
    """
    # Increment time step
    self.t += 1

    # Bias correction factors (computed once per update)

```

```

bias_correction1 = 1 - self.beta1 ** self.t
bias_correction2 = 1 - self.beta2 ** self.t

# Update parameters for all layers
for l in range(1, self.num_layers + 1):
    grad_W = gradients[l]['W']
    grad_b = gradients[l]['b']

    # =====
    # Update W^(l) (weight matrix for layer l)
    # =====
    # First moment: m = β1·m + (1-β1)·g
    self.m_W[l] = self.beta1 * self.m_W[l] + (1 - self.beta1) * grad_W

    # Second moment: v = β2·v + (1-β2)·g2
    self.v_W[l] = self.beta2 * self.v_W[l] + (1 - self.beta2) * (grad_W ** 2)

    # Bias-corrected moments
    m_hat_W = self.m_W[l] / bias_correction1
    v_hat_W = self.v_W[l] / bias_correction2

    # Parameter update: θ = θ - η·m/(√v + ε)
    self.W[l] -= self.lr * m_hat_W / (np.sqrt(v_hat_W) + self.epsilon)

    # =====
    # Update b^(l) (bias vector for layer l)
    # =====
    self.m_b[l] = self.beta1 * self.m_b[l] + (1 - self.beta1) * grad_b
    self.v_b[l] = self.beta2 * self.v_b[l] + (1 - self.beta2) * (grad_b ** 2)

    m_hat_b = self.m_b[l] / bias_correction1
    v_hat_b = self.v_b[l] / bias_correction2

    self.b[l] -= self.lr * m_hat_b / (np.sqrt(v_hat_b) + self.epsilon)

def train(self, X_train: np.ndarray, y_train: np.ndarray,

```

```
X_val: np.ndarray = None, y_val: np.ndarray = None,  
epochs: int = 500000, print_every: int = 10000,  
verbose: bool = True) → Dict[str, Dict[int, float]]:  
"""
```

Training loop with validation monitoring

Algorithm:

For each epoch:

1. Forward pass: compute predictions
2. Compute loss: measure error
3. Backward pass: compute gradients
4. Update parameters: apply Adam optimizer
5. (Optional) Validate and log progress

Args:

```
X_train: Training inputs of shape (n_train, p)  
y_train: Training labels of shape (n_train,)  
X_val: Validation inputs of shape (n_val, p) [optional]  
y_val: Validation labels of shape (n_val,) [optional]  
epochs: Number of training iterations  
print_every: Frequency of progress logging  
verbose: Whether to print progress
```

Returns:

history: Dictionary containing training and validation losses

```
"""
```

```
history = {  
    'train_loss': {},  
    'val_loss': {}  
}
```

for epoch in range(epochs):

```
# ======  
# Training Step  
# ======  
# 1. Forward pass
```

```

y_train_pred = self.forward(X_train)

# 2. Compute loss
train_loss = self.loss_fn(y_train, y_train_pred)
history['train_loss'][epoch] = train_loss

# 3. Backward pass
gradients = self.backward(y_train)

# 4. Update parameters
self.update_parameters(gradients)

# =====
# Validation Step (if validation data provided)
# =====
if X_val is not None and y_val is not None:
    y_val_pred = self.forward(X_val)
    val_loss = self.loss_fn(y_val, y_val_pred)
    history['val_loss'][epoch] = val_loss
else:
    val_loss = None

# =====
# Logging
# =====
if verbose and (epoch % print_every == 0 or epoch == epochs - 1):
    log_str = f"Epoch {epoch:07d}/{epochs:07d} | Train Loss: {train_los
s:.5f}"
    if val_loss is not None:
        log_str += f" | Val Loss: {val_loss:.5f}"

    # Compute accuracy
    train_acc = np.mean((y_train_pred > 0.5) == y_train)
    val_acc = np.mean((y_val_pred > 0.5) == y_val)
    log_str += f" | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f}"

```

```
    print(log_str)

    return history

def predict(self, X: np.ndarray, threshold: float = 0.5) → np.ndarray:
    """
    Make binary predictions
```

Mathematical operation:

1. prob = forward(X) → compute $P(y=1 | x)$
2. prediction = $\mathbb{1}(\text{prob} > \text{threshold})$

Args:

X: Input data of shape (n, p)
threshold: Decision threshold (default 0.5)

Returns:

predictions: Binary predictions of shape (n,)

"""

```
prob = self.forward(X)
predictions = (prob > threshold).astype(int)
return predictions
```

```
def predict_proba(self, X: np.ndarray) → np.ndarray:
```

"""

Predict probabilities $P(y=1 | x)$

Args:

X: Input data of shape (n, p)

Returns:

probabilities: Probabilities of shape (n,)

"""

```
return self.forward(X)
```