



STATS-413 HW-4

Author: Hochan Son

UID: 206547205

Date: @November 6, 2025

Problem 1 “Barack Obama” example. Consider a MLP with one hidden layer. Let x be the input (Barack), h be the hidden layer, and y be the output (Obama). Suppose both x and y are words in the same dictionary, where x is the current word, and y is the next word. Both x and y are on-hot vector.

$$\begin{aligned} & \text{Let } h = W_{\text{embed}}x, \\ & s = W_{\text{unembed}}h, \\ & p = \text{softmax}(s), \\ & y \sim p \\ & p(y_c = 1|s) = p_c \end{aligned}$$

where

- $x = \text{'Barack'}$: input layer
- $h = W_{\text{embed}}$: hidden layer
- $s = W_{\text{unembed}} h$:
- $p = \text{softmax}(s)$: probability of the output distribution
- $y = \text{'Obama'}$: predictor

(1) What are the dimensionalities of W_{embed} and W_{unembed} ?

Interpret the meaning of the columns of W_{embed}

Let's say:

- Word size = $W(100k)$ and,
- Hidden layer size = H (100 dimensions)

The dimensionality of each layers

- $W_{embed} = \text{shape}(H \times W)$
 - W_K is a dimension of one-hot embedding on K -th elements
 - H is hidden layer representation
- $W_{unembed} = \text{shape}(W \times H)$
 - H is a hidden states ⊙ W -dimension logits score(for each word)

(2) Let $J = \log p(y|s)$, show that $\partial J / \partial s = y - p = e$.

Calculate $\partial J / \partial h$, $\partial J / \partial W_{unembed}$, and $\partial J / \partial W_{embed}$ with chain rule
In your calculation, you can first pretend all the vectors and matrices are scalars (one-dimensional numbers), and then guess the forms of the general results.

Part 1: Loss Function

starting with softmax cross-entropy:

$$J = \log p(y|s)$$

$$\begin{aligned} J &= \log p(y|s) \\ &= \sum_c y_c \log p_c \end{aligned}$$

For one-hot encoded target y (only one correct class):

$$= \sum_c y_c \log p_c$$

Since y is one-hot, only one term survives (where $y_c = 1$):

$$\begin{aligned} ; p(y|s) &= \prod_c p_c^{y_c} \\ p_c &= \frac{e^{s_c}}{z} \text{ where } \sum_c e^{s_c} \\ J &= \sum_c y_c (s_c - \log z) \\ &= \sum_c y_c s_c - \log z \sum_c y_c \\ &= (\sum_c y_c s_c) - \log z \end{aligned}$$

Since $\sum y_c = 1$:

$$J = \sum_c y_c s_c - \log z$$

Part 2: Gradient w.r.t Logits ($\partial J / \partial s$)

$$\frac{\partial J}{\partial s_c} = y_c - \frac{1}{z} \frac{\partial z}{\partial s_c}$$

Since $z = \sum e^s$:

$$\frac{\partial z}{\partial s} = e^s$$

Therefore:

$$\frac{\partial J}{\partial s} = y - p = \text{error}$$

Part 3: Gradient w.r.t. Hidden State ($\partial J / \partial h$) :

using the chain rule:

$$\frac{\partial J}{\partial h_i} = \sum_k \frac{\partial J}{\partial s_k} \cdot \frac{\partial s_k}{\partial h_i}$$

Since $s_k = \sum_j W_{kj}^u h_j$:

$$\frac{\partial s_k}{\partial h_i} = W_{ki}^u$$

Therefore:

$$\frac{\partial J}{\partial h_i} = \sum_k \frac{\partial J}{\partial s_k} \cdot W_{ki}^u$$

In matrix form:

$$\frac{\partial J}{\partial h} = W_{unembed}^T \cdot \frac{\partial J}{\partial s}$$

Part 4: Gradient w.r.t. Unembed Weights ($\partial J / \partial W_{unembed}$):

$$\frac{\partial J}{\partial W_{jk}^u} = \frac{\partial J}{\partial s_k} \cdot \frac{\partial s_k}{\partial W_{kj}^u}$$

Since $s_k = \sum_j W_{kj}^u h_j$

$$\frac{\partial s_k}{\partial W_{kj}^u} = h_j$$

Therefore:

$$\boxed{\frac{\partial J}{\partial W_{unembed}} = \frac{\partial J}{\partial s} \cdot h^T}$$

Part 5: Gradient w.r.t Embed Weights ($\partial J / \partial W_{embed}$):

$$\frac{\partial J}{\partial W_{kj}^e} = \frac{\partial J}{\partial h_i} \cdot \frac{\partial h_i}{\partial W_{kj}^e}$$

Since $h_i = \sum_k W_{ik}^e x_k$:

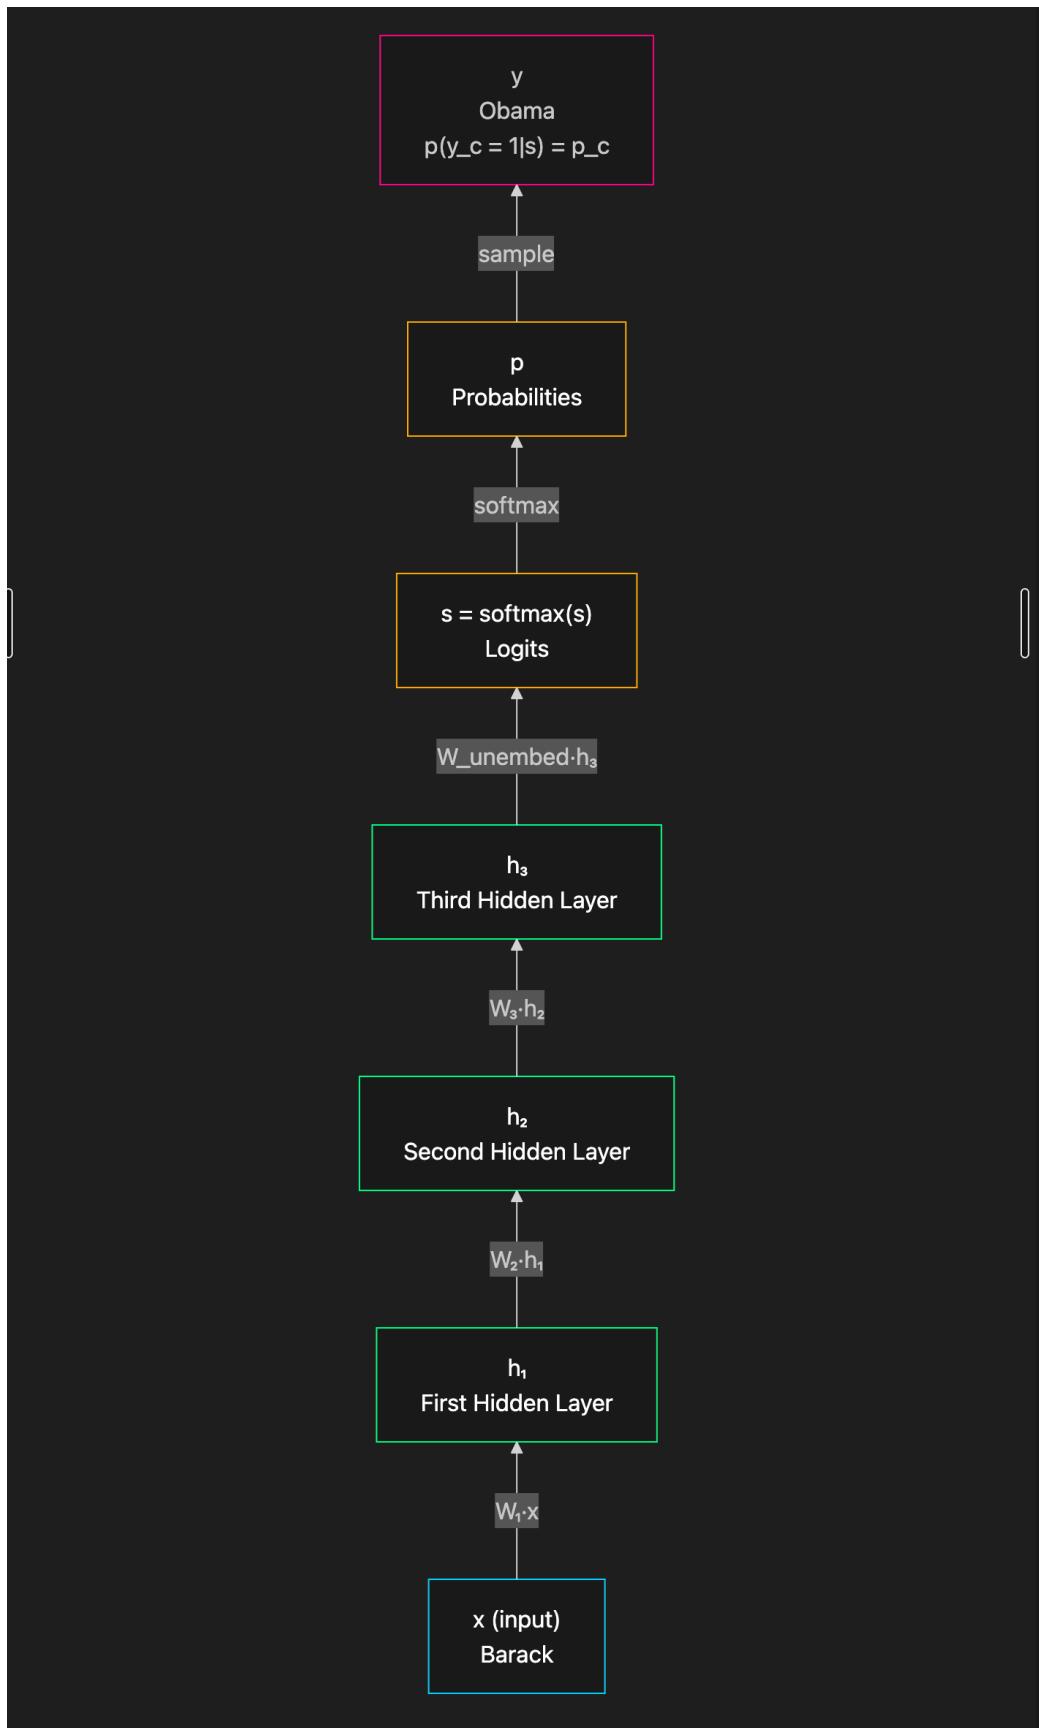
$$\frac{\partial h_i}{\partial W_{kj}^e} = x_j$$

Therefore:

$$\boxed{\frac{\partial J}{\partial W_{embed}} = \frac{\partial J}{\partial h} \cdot x^T}$$

(3) Draw a diagram of network with multiple layers of latent vectors

The multi layer perceptron:



```

graph BT
x["x (input)<br/>Barack"] -->|W1·x| h1["h1<br/>First Hidden Layer"]
h1 -->|W2·h1| h2["h2<br/>Second Hidden Layer"]
h2 -->|W3·h2| h3["h3<br/>Third Hidden Layer"]
h3 -->|Wunembed·h3| s["s = softmax(s)<br/>Logits"]
s -->|softmax| p["p<br/>Probabilities"]
p -->|sample| y["y<br/>Obama<br/>p(yc = 1|s) = pc"]

style x fill:#1a1a1a,stroke:#00d4ff,color:#fff
style h1 fill:#1a1a1a,stroke:#00ff88,color:#fff
style h2 fill:#1a1a1a,stroke:#00ff88,color:#fff
style h3 fill:#1a1a1a,stroke:#00ff88,color:#fff
style s fill:#1a1a1a,stroke:#ffaa00,color:#fff
style p fill:#1a1a1a,stroke:#ffaa00,color:#fff
style y fill:#1a1a1a,stroke:#ff0088,color:#fff

```

Where:

- Input (x): “Barack” as one-hot vector
- 3 hidden layers (h₁,h₂,h₃): Each layer transforms the previous representation through learned weight matrices
- Logits (s): Final hidden layer projected back to vocabulary space
- Probability (p): Softmax applied to logits
- Output (y): Sampled next word “Obama”

Problem 2. “I love machine learning” example. Suppose we observe “I love machine” and we want to predict the next word.

Let x_t be one-hot vectors,

$$x_1 = "I",$$

$$x_2 = "love",$$

$$x_3 = "machine".$$

Let h_t be the hidden vectors,

$$h_0 = 0,$$

$$h_1 = \tanh(W_{embed}X_1 + W_{recurrent}h_0),$$

$$h_2 = \tanh(W_{embed}X_2 + W_{recurrent}h_1),$$

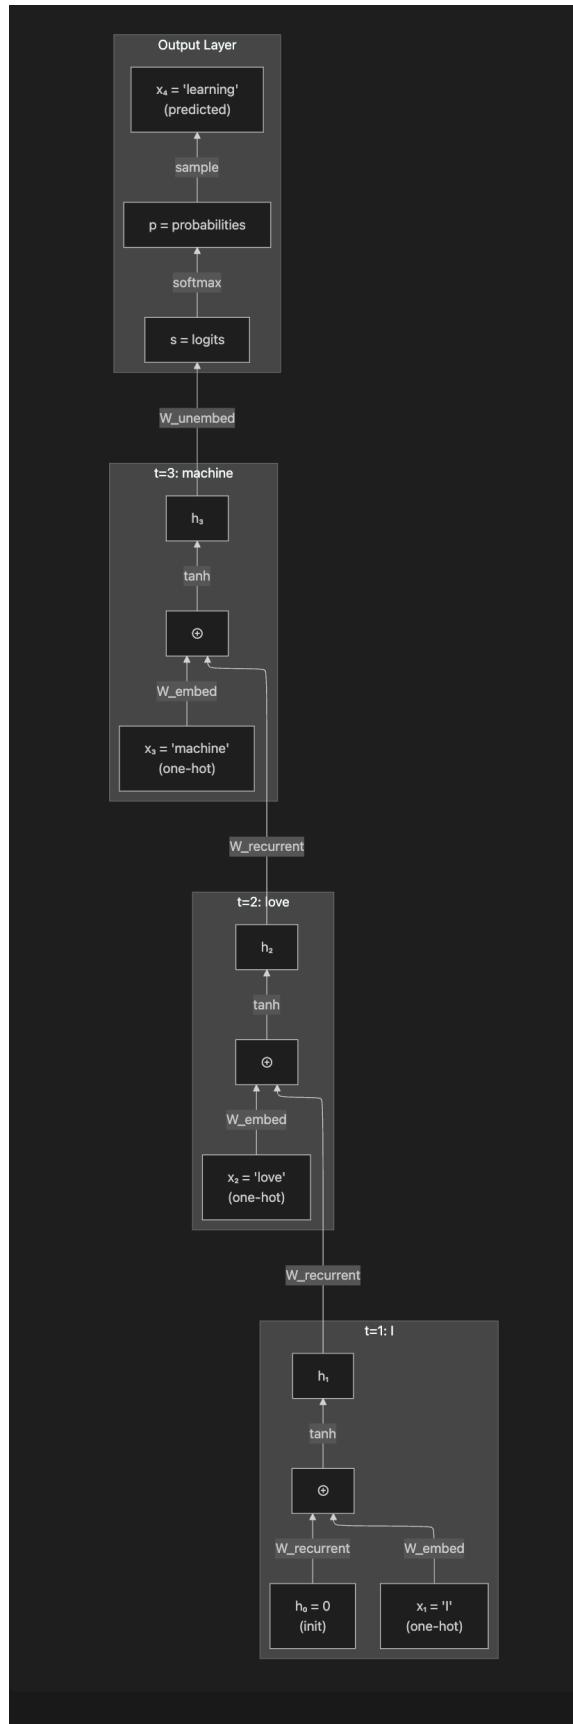
$$h_3 = \tanh(W_{embed}X_3 + W_{recurrent}h_2), \text{ and}$$

$$s = W_{unembed}h_3,$$

$$p = softmax(s),$$

x_4 is sampled according to p .

(1) Draw a diagram to illustrate the model



```

graph BT
    subgraph t1["t=1: I"]
        x1["x1 = 'I'<br/>(one-hot)"]
        h0["h0 = 0<br/>(init)"]
        x1 --|W_embed|--> add1["⊕"]
        h0 --|W_recurrent|--> add1
        add1 --|tanh|--> h1["h1"]
    end

    subgraph t2["t=2: love"]
        x2["x2 = 'love'<br/>(one-hot)"]
        x2 --|W_embed|--> add2["⊕"]
        h1 --|W_recurrent|--> add2
        add2 --|tanh|--> h2["h2"]
    end

    subgraph t3["t=3: machine"]
        x3["x3 = 'machine'<br/>(one-hot)"]
        x3 --|W_embed|--> add3["⊕"]
        h2 --|W_recurrent|--> add3
        add3 --|tanh|--> h3["h3"]
    end

    subgraph output["Output Layer"]
        h3 --|W_unembed|--> s["s = logits"]
        s --|softmax|--> p["p = probabilities"]
        p --|sample|--> x4["x4 = 'learning'<br/>(predicted)"]
    end

```

(2) Let $J = \log p(x_4 | s)$.

Calculate $\partial J / \partial W_{embed}$, $\partial J / \partial W_{unembed}$, and $\partial J / \partial W_{recurrent}$. **In your calculation, you can first pretend all the vectors and matrices are scalars (one-dimensional numbers), and then guess the forms of the general results.**

Let x_t be one-hot vectors. $x_1 = "I"$, $x_2 = "love"$, $x_3 = "machine"$

Let $h_0 = 0$, $h_t = \tanh(z_1)$ for $t = 1, 2, 3$ where $z_1 = W_{embed}x_t + W_{recurrent}h_{t-1}$

Let $s = W_{unembed}h_3$, $p = softmax(s)$, $x_4 \sim p$

Step 1. Forward pass:

$$\begin{aligned}
h_0 &= 0 \\
z_1 &= W_{embed}x_1, h_1 = \tanh(z_1) \\
z_1 &= W_{embed}x_2 + W_{recurrent}h_1, h_2 = \tanh(z_2) \\
z_1 &= W_{embed}x_3 + W_{recurrent}h_2, h_3 = \tanh(z_3) \\
s &= W_{unembed}h_3, p = \text{softmax}(s), J = \log p(x_4|s)
\end{aligned}$$

Key Gradients (given earlier):

$$\frac{\partial J}{\partial s} = y - p = e \text{ (error signal)}$$

Define the backward errors:

$$\delta_t = \frac{\partial J}{\partial h_t}$$

Back propagate the error:

Recall: $s = W_{unembed}h_3$, and $\partial J / \partial s = e = (y - p)$

$$\begin{aligned}
\delta_3 &= \frac{\partial J}{\partial h_3} = \frac{\partial J}{\partial s} \cdot \frac{\partial s}{\partial h_3} = e \cdot W_{unembed}^T \\
\delta_2 &= \frac{\partial J}{\partial h_2} = \frac{\partial J}{\partial h_3} \cdot \frac{\partial h_3}{\partial h_2} = \delta_3 \cdot \sigma'(z_3) \cdot W_{recurrent}^T \\
\delta_1 &= \frac{\partial J}{\partial h_1} = \frac{\partial J}{\partial h_2} \cdot \frac{\partial h_2}{\partial h_1} = \delta_2 \cdot \sigma'(z_2) \cdot W_{recurrent}^T
\end{aligned}$$

Let trace the error propagates all the way back to h_1 :

$$\delta_1 = \delta_2 \cdot \sigma'(z_2) \cdot W_{recurrent}^T$$

Substitute δ_2 :

$$\delta_1 = [\delta_3 \cdot \sigma'(z_3) \cdot W_{recurrent}^T] \cdot \sigma'(z_2) \cdot W_{recurrent}^T$$

Substitute δ_3 :

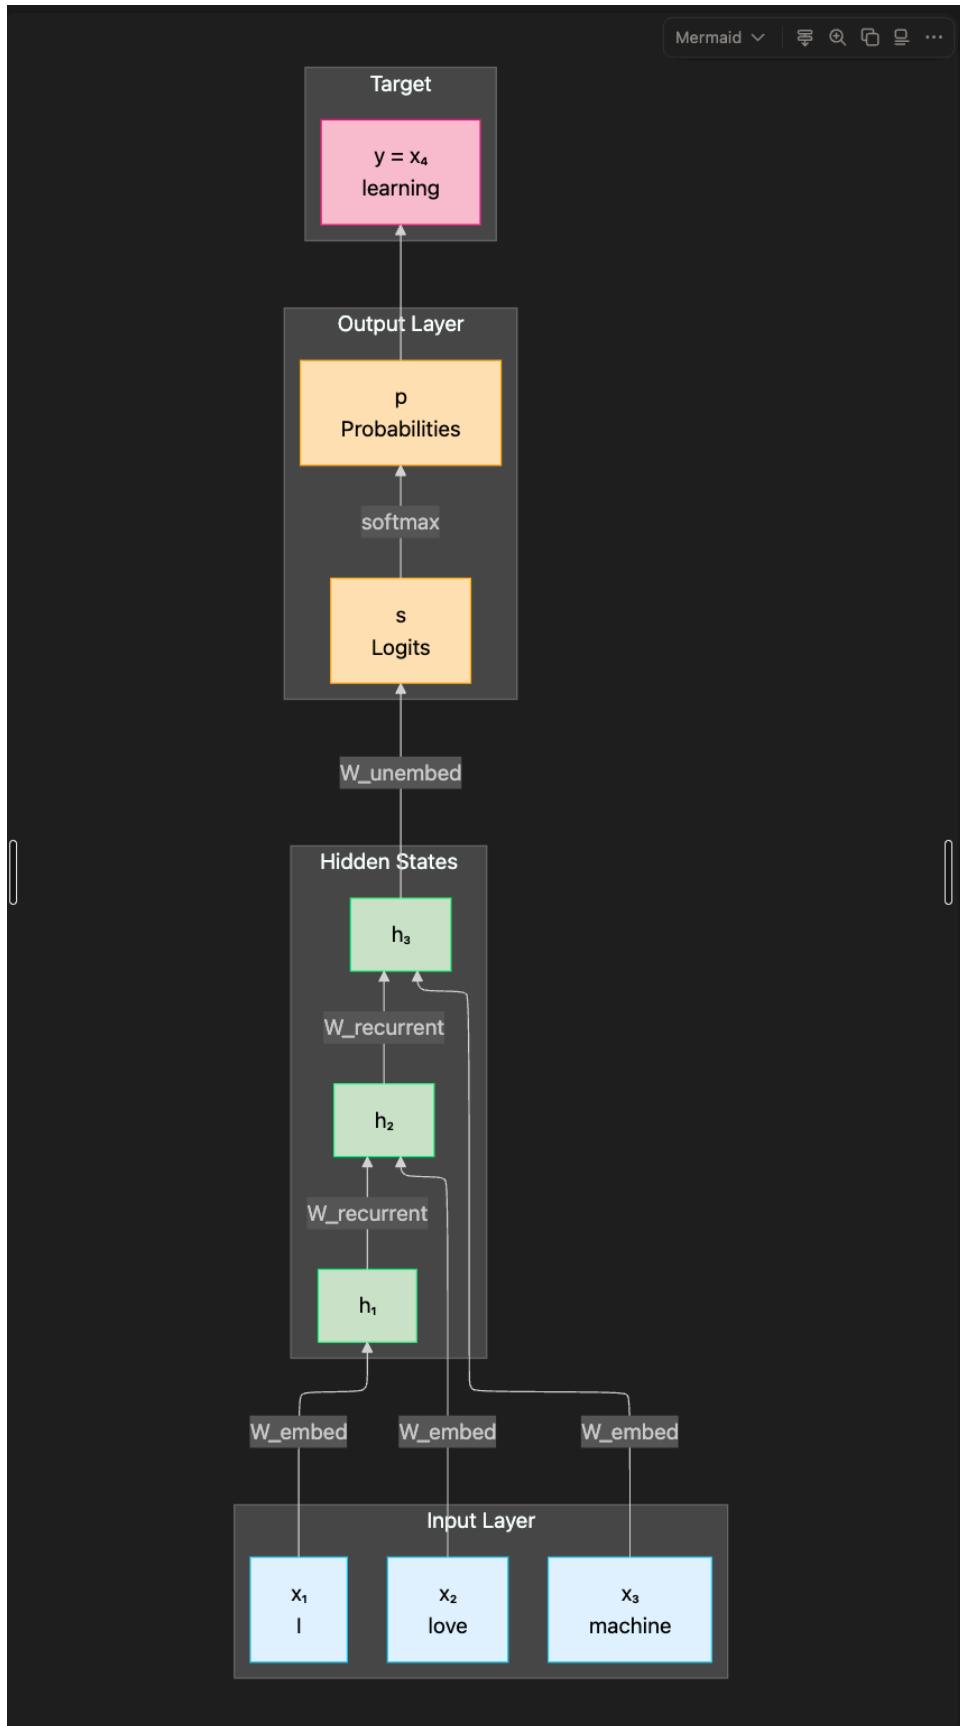
$$\begin{aligned}
\delta_1 &= [W_{recurrent}^T \cdot e] \cdot \sigma'(z_3) \cdot W_{recurrent}^T \cdot \sigma'(z_2) \cdot W_{recurrent}^T \\
&= W_{recurrent}^T \cdot e \cdot \sigma'(z_3) \cdot W_{recurrent}^T \cdot \sigma'(z_2) \cdot W_{recurrent}^T
\end{aligned}$$

where $\sigma'(z_t) = 1 - \tanh^2(z_t)$

- Calculate $\partial J / \partial W_{unembed}$

$$\begin{aligned}
\frac{\partial J}{\partial W_{unembed}} &= \frac{s = W_{unembed}h_3}{\frac{\partial J}{\partial s} \cdot \frac{\partial s}{\partial W_{unembed}}} \\
&= \boxed{\frac{\partial J}{\partial W_{unembed}} = e \cdot h_3}
\end{aligned}$$

(3) Draw a diagram of network with multiple recurrent layers of latent vector



```

graph BT
    subgraph input["Input Layer"]
        x1["x1<br/>I"]
        x2["x2<br/>love"]
        x3["x3<br/>machine"]
    end

    subgraph hidden["Hidden States"]
        h1["h1"]
        h2["h2"]
        h3["h3"]
    end

    subgraph output["Output Layer"]
        s["s<br/>Logits"]
        p["p<br/>Probabilities"]
    end

    subgraph target["Target"]
        y["y = x4<br/>learning"]
    end

    x1 → |W_embed| h1
    x2 → |W_embed| h2
    x3 → |W_embed| h3

    h1 → |W_recurrent| h2
    h2 → |W_recurrent| h3

    h3 → |W_unembed| s
    s → |softmax| p
    p → y

    style x1 fill:#e1f5ff,stroke:#00d4ff,color:#000
    style x2 fill:#e1f5ff,stroke:#00d4ff,color:#000
    style x3 fill:#e1f5ff,stroke:#00d4ff,color:#000
    style h1 fill:#c8e6c9,stroke:#00ff88,color:#000
    style h2 fill:#c8e6c9,stroke:#00ff88,color:#000
    style h3 fill:#c8e6c9,stroke:#00ff88,color:#000
    style s fill:#ffe0b2,stroke:#ffaa00,color:#000
    style p fill:#ffe0b2,stroke:#ffaa00,color:#000
    style y fill:#f8bbd0,stroke:#ff0088,color:#000

```

$$h_0 \rightarrow h_1 \odot W_{embed} x_1 \rightarrow h_2 \odot W_{embed} x_2 \rightarrow h_3 \odot W_{embed} x_3 \rightarrow s : W_{unembed} h_3 \rightarrow p \rightarrow x_4$$

Problem 3 Residual steam, For the “I love machine learning” example, consider the residual parameterization:

- $h_t = h_{t-1} + \tanh(W_{recurrent} h_{t-1} + W_{embed} x_t)$
- Let $J = \log p(x_4|s)$, where $s = W_{unembed} h_3$,

Starting from $\partial J / \partial h_3$, calculate $\partial J / \partial h_1$.

Step 1: Forward Pass (computing hidden states)

At t=1:

$$\begin{aligned} z_1 &= W_{recurrent} h_0 + W_{embed} x_1 \\ h_1 &= h_0 + \tanh(z_1) = 0 + \tanh(z_1) = \tanh(z_1) \end{aligned}$$

At t=2:

$$\begin{aligned} z_2 &= W_{recurrent} h_1 + W_{embed} x_2 \\ h_2 &= h_1 + \tanh(z_2) \end{aligned}$$

At t=3:

$$\begin{aligned} z_3 &= W_{recurrent} h_2 + W_{embed} x_3 \\ h_3 &= h_2 + \tanh(z_3) \end{aligned}$$

Step 2: Backward pass(Gradient)

Given $\partial J / \partial h_3$

the gradient from the output layer:

$$\frac{\partial J}{\partial h_3} = (\text{Given from loss computation})$$

From earlier problems we know:

$$\frac{\partial J}{\partial s} = y - p = e(\text{error})$$

where $s = W_{recurrent} h_3$ and $p = \text{softmax}(s)$

Therefore,

$$\frac{\partial J}{\partial h_3} = \frac{\partial J}{\partial s} \cdot \frac{\partial s}{\partial h_3} = e \cdot W_{unembed}^T$$

Back propagation h3→h2

Step 1: Set up the Jacobian:

Recall the residual update equation:

$$h_3 = h_2 + \tanh(z_3)$$

where $z_3 = W_{recurrent}h_2 + W_{embed}x_3$

Taking the derivative w.r.t. h_2 :

$$\frac{\partial h_3}{\partial h_2} = \frac{\partial}{\partial h_2}(h_2 + \tanh(z_3))$$

split into two terms,

$$\begin{aligned} &= \frac{\partial h_2}{\partial h_2} + \frac{\partial \tanh(z_3)}{\partial h_2} \\ &= I + \frac{\partial \tanh(z_3)}{\partial h_2} \end{aligned}$$

Step 2: Apply the chain rule to the tanh term

For the tanh derivative, use the chain rule:

$$\frac{\partial \tanh(z_3)}{\partial h_2} = \frac{\partial \tanh(z_3)}{\partial z_3} \times \frac{\partial z_3}{\partial h_2}$$

1. Compute tanh derivative rule:

$$\frac{d}{du} \tanh(u) = 1 - \tanh^2(u)$$

so the tanh derivative is:

$$\frac{\partial \tanh(z_3)}{\partial z_3} = 1 - \tanh^2(z_3) = \sigma'(z_3)$$

2. Compute $\frac{\partial z_3}{\partial h_2}$

Recall:

$$z_3 = W_{recurrent}h_2 + W_{embed}x_3$$

Taking the derivative w.r.t. h_2 :

$$\frac{\partial z_3}{\partial h_2} = \frac{\partial}{\partial h_2}(W_{recurrent}h_2 + W_{embed}x_3)$$

The $W_{embed}x_3$ term doesn't depend on h_2 , so its derivative is 0.

$$= \frac{\partial}{\partial h_2}(W_{recurrent}h_2) = W_{recurrent} : \text{weight matrix}$$

3. Multiply the two derivatives (chain rule)

$$\begin{aligned}
\frac{\partial \tanh(z_3)}{\partial h_2} &= \frac{\partial \tanh(z_3)}{\partial z_3} \cdot \frac{\partial z_3}{\partial h_2} \\
&= (1 - \tanh^2(z_3)) \cdot W_{recurrent} \\
&= \sigma'(z_3) \cdot W_{recurrent}
\end{aligned}$$

Therefore:

$$\boxed{\frac{\partial h_3}{\partial h_2} = I + \sigma'(z_3) \cdot W_{recurrent}}$$

where $\sigma'(z_3) = 1 - \tanh^2(z_3)$ is a scalar (or diagonal matrix when considering batch dimensions)

4. apply the chain rule for gradients

$$\frac{\partial J}{\partial h_2} = \frac{\partial J}{\partial h_3} \cdot \frac{\partial h_3}{\partial h_2}$$

Substitute the Jacobian:

$$\boxed{\frac{\partial J}{\partial h_2} = \frac{\partial J}{\partial h_3} (I + \sigma'(z_3) W_{recurrent})}$$

5. Expand element-wise distribute the multiplication

$$\frac{\partial J}{\partial h_2} = \frac{\partial J}{\partial h_3} I + \frac{\partial J}{\partial h_3} \sigma'(z_3) W_{recurrent}$$

Simplify (since $x \cdot I = x$):

$$\boxed{\frac{\partial J}{\partial h_2} = \underbrace{\frac{\partial J}{\partial h_3}}_{direct\ path} + \underbrace{\frac{\partial J}{\partial h_3} \sigma'(z_3) W_{recurrent}}_{recurrent\ path}}$$

Observations:

Two gradient paths:

1. Direct path (Identity term):

- Gradient flows unchanged through I (identity)
- No attenuation

2. Recurrent path (Tanh term):

- Gradient multiplied by $\sigma'(z_3) \subset (0, 1)$ and $W_{recurrent}$
- May vanish

Even if the recurrent path vanishes, the direct path keeps gradient flowing!

Problem 4 Please play with the PyTorch code provided by the following webpage:

<https://machinelearningmastery.com/text-generation-with-lstm-in-pytorch/>

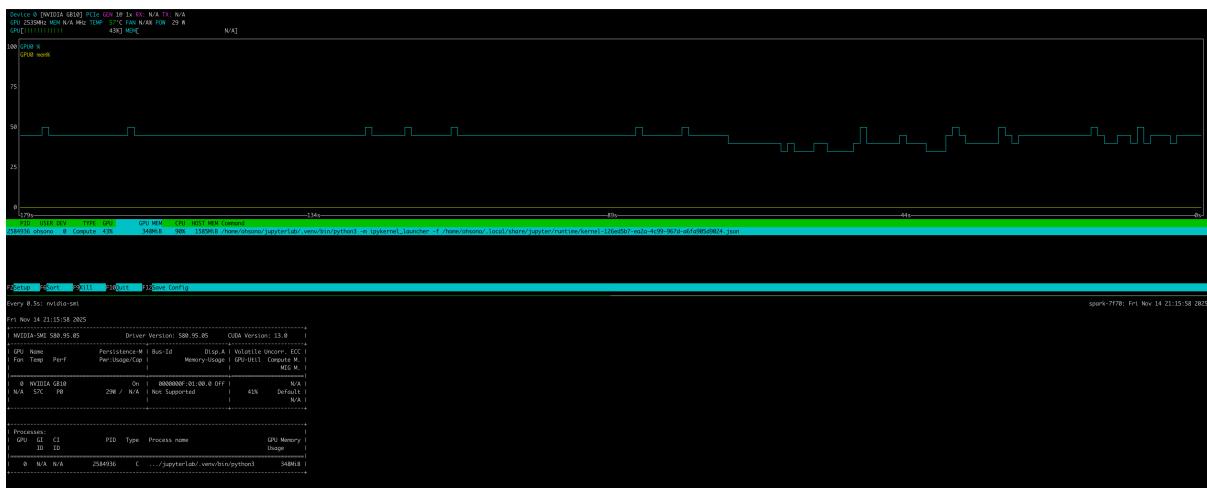
Please write a brief explanation of the code and show your results. You can explore the code by varying the design parameters.

What I used a text for training LSTM model

- The prince, Nicolo Machiavelli
 - Source, <https://www.gutenberg.org/files/1232/1232-h/1232-h.htm>
 - Parameters:

Using device: cuda
Loaded 282277 characters
Vocab size: 55
Dataset size: 282277 sequences
Model parameters: 942775

- Machine Specs:
 - DGX-Spark, GB10, 128GB Memory
 - GPU utilization : 43% / Memory Utilization: 16GB



- CODE SAMPLE

COMPLETE PyTorch LSTM Text Generation Implementation

Based on MachineLearningMastery approach

```
import torch  
import torch.nn as nn
```

```

import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import numpy as np

# =====
===
# STEP 1: DATA PREPARATION
# =====
===

class TextDataset(Dataset):
    """Convert text into character sequences for LSTM training"""

    def __init__(self, text, seq_length=50):
        """
        Args:
            text: Raw text string
            seq_length: Length of input sequences
        """
        self.seq_length = seq_length

        # Build character vocabulary
        self.chars = sorted(set(text))
        self.char_to_idx = {c: i for i, c in enumerate(self.chars)}
        self.idx_to_char = {i: c for i, c in enumerate(self.chars)}

        # Encode entire text to indices
        self.text_encoded = [self.char_to_idx[c] for c in text]

    def __len__(self):
        return len(self.text_encoded) - self.seq_length

    def __getitem__(self, idx):
        """Return (input_sequence, target_sequence) pair"""
        x = torch.tensor(
            self.text_encoded[idx:idx + self.seq_length],
            dtype=torch.long
        )
        y = torch.tensor(
            self.text_encoded[idx + 1:idx + self.seq_length + 1],
            dtype=torch.long
        )
        return x, y

    def decode(self, indices):
        """Convert indices back to text"""
        return ''.join([self.idx_to_char[i] for i in indices])

```

```

# =====
===
# STEP 2: MODEL ARCHITECTURE
# =====
===

class LSTMTextGenerator(nn.Module):
    """LSTM-based character-level text generation model"""

    def __init__(self, vocab_size, embedding_dim=128, hidden_dim=256,
                 num_layers=2, dropout=0.5):
        """
        Args:
            vocab_size: Number of unique characters
            embedding_dim: Dimension of character embeddings
            hidden_dim: Dimension of LSTM hidden state
            num_layers: Number of stacked LSTM layers
            dropout: Dropout rate between LSTM layers
        """
        super(LSTMTextGenerator, self).__init__()

        self.embedding = nn.Embedding(vocab_size, embedding_dim)

        self.lstm = nn.LSTM(
            input_size=embedding_dim,
            hidden_size=hidden_dim,
            num_layers=num_layers,
            dropout=dropout if num_layers > 1 else 0.0,
            batch_first=True
        )

        self.fc = nn.Linear(hidden_dim, vocab_size)
        self.hidden_dim = hidden_dim
        self.num_layers = num_layers

    def forward(self, x, hidden=None):
        """
        Args:
            x: Input tensor of shape (batch_size, seq_length)
            hidden: Tuple of (hidden_state, cell_state) or None
        Returns:
            logits: Output logits of shape (batch_size, seq_length, vocab_size)
            hidden: Updated (hidden_state, cell_state)
        """

```

```

# Embedding: (batch, seq_len) → (batch, seq_len, embed_dim)
embedded = self.embedding(x)

# LSTM: (batch, seq_len, embed_dim) → (batch, seq_len, hidden_dim)
lstm_out, hidden = self.lstm(embedded, hidden)

# Linear: (batch, seq_len, hidden_dim) → (batch, seq_len, vocab_size)
logits = self.fc(lstm_out)

return logits, hidden

# =====
===
# STEP 3: TRAINING LOOP
# =====
===

def train_epoch(model, train_loader, criterion, optimizer, device):
    """Train for one epoch"""
    model.train()
    total_loss = 0

    for batch_idx, (x, y) in enumerate(train_loader):
        x, y = x.to(device), y.to(device)

        # Forward pass
        logits, _ = model(x) # (batch, seq_len, vocab_size)

        # Compute loss: reshape for CrossEntropyLoss
        # CrossEntropyLoss expects (N, C) where N = batch*seq_len, C = vocab_size
        loss = criterion(
            logits.view(-1, logits.size(-1)), # (batch*seq_len, vocab_size)
            y.view(-1) # (batch*seq_len,)
        )

        # Backward pass
        optimizer.zero_grad()
        loss.backward()

        # Optional: gradient clipping to prevent exploding gradients
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)

        optimizer.step()
        total_loss += loss.item()

    avg_loss = total_loss / len(train_loader)

```

```

return avg_loss

def train(model, train_loader, num_epochs, learning_rate, device='cpu'):
    """Train the LSTM model"""
    model.to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    for epoch in range(num_epochs):
        loss = train_epoch(model, train_loader, criterion, optimizer, device)
        print(f"Epoch {epoch+1}/{num_epochs}, Loss: {loss:.4f}")

# =====
# STEP 4: TEXT GENERATION
# =====

def generate_text(model, dataset, seed_text, length, temperature=1.0,
                  device='cpu'):
    """
    Generate text starting from seed_text

    Args:
        model: Trained LSTM model
        dataset: TextDataset instance (for encoding/decoding)
        seed_text: Starting text
        length: Number of characters to generate
        temperature: Controls randomness
            - Low (0.5): More deterministic
            - High (1.5): More random
        device: 'cpu' or 'cuda'

    Returns:
        Generated text string
    """
    model.eval()

    # Convert seed to indices
    indices = [dataset.char_to_idx[c] for c in seed_text]

    with torch.no_grad():
        for _ in range(length):
            # Use last seq_length characters as context

```

```

if len(indices) >= dataset.seq_length:
    x = torch.tensor(
        indices[-dataset.seq_length:],
        dtype=torch.long
    ).unsqueeze(0).to(device)
else:
    x = torch.tensor(
        indices,
        dtype=torch.long
    ).unsqueeze(0).to(device)

# Get model prediction
logits, _ = model(x)

# Get logits for next character (last position in sequence)
next_logits = logits[0, -1, :] / temperature

# Apply softmax and sample
probs = torch.softmax(next_logits, dim=0).cpu().numpy()
next_idx = np.random.choice(len(probs), p=probs)

indices.append(next_idx)

return dataset.decode(indices)

# =====
===
# STEP 5: COMPLETE USAGE EXAMPLE
# =====
===

if __name__ == "__main__":
    # Configuration
    TEXT_FILE = "theprince.txt" # Your text file
    SEQ_LENGTH = 50
    BATCH_SIZE = 32
    EMBEDDING_DIM = 128
    HIDDEN_DIM = 256
    NUM_LAYERS = 2
    DROPOUT = 0.3
    LEARNING_RATE = 0.001
    NUM_EPOCHS = 50
    DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    print(f"Using device: {DEVICE}")

```

```

# ===== Load and Prepare Data =====
# Load your text file
with open(TEXT_FILE, 'r', encoding='utf-8') as f:
    text = f.read().lower() # Lowercase for consistency

print(f"Loaded {len(text)} characters")

# Create dataset and dataloader
dataset = TextDataset(text, seq_length=SEQ_LENGTH)
train_loader = DataLoader(
    dataset,
    batch_size=BATCH_SIZE,
    shuffle=True,
    pin_memory=True if DEVICE.type == 'cuda' else False
)

print(f"Vocab size: {len(dataset.chars)}")
print(f"Dataset size: {len(dataset)} sequences")

# ===== Create Model =====
model = LSTMTextGenerator(
    vocab_size=len(dataset.chars),
    embedding_dim=EMBEDDING_DIM,
    hidden_dim=HIDDEN_DIM,
    num_layers=NUM_LAYERS,
    dropout=DROPOUT
)

print(f"Model parameters: {sum(p.numel() for p in model.parameters())}")

# ===== Train =====
train(
    model,
    train_loader,
    num_epochs=NUM_EPOCHS,
    learning_rate=LEARNING_RATE,
    device=DEVICE
)

# ===== Generate Text =====
seed = "the great"
print(f"\nGenerating text starting with: '{seed}'")

for temperature in [0.5, 1.0, 1.5]:
    print(f"\nTemperature: {temperature}")
    generated = generate_text(
        model,

```

```

        dataset,
        seed,
        length=200,
        temperature=temperature,
        device=DEVICE
    )
    print(generated)

# ===== Save Model =====
torch.save(model.state_dict(), 'lstm_model.pt')
print("\nModel saved to 'lstm_model.pt'")

```

- Evaluation

```

"""
LSTM TEXT GENERATION EVALUATION TOOLKIT
Compute metrics to analyze generated text quality
"""

import numpy as np
from collections import Counter
from typing import Dict, List, Tuple
import math

# =====
# BASIC TEXT METRICS
# =====
# =====

class TextMetrics:
    """Compute various metrics on text"""

    @staticmethod
    def distinctness(text: str) -> float:
        """
        Compute distinctness (Type-Token Ratio)

        Range: 0-1
        Higher = more diverse vocabulary
        Lower = more repetitive

        Args:
            text: Generated text

        Returns:
    
```

```

    Distinctness score
    """
words = text.lower().split()
if len(words) == 0:
    return 0.0

unique_words = len(set(words))
total_words = len(words)

return unique_words / total_words

@staticmethod
def repetition_ratio(text: str) → float:
    """
Count adjacent repeated words

Range: 0-1
Higher = more repetitive (bad)
Lower = more diverse (good)

Example:
"the the the and the" → repetition_ratio = 0.75
"""
words = text.lower().split()
if len(words) < 2:
    return 0.0

repetitions = sum(1 for i in range(len(words) - 1) if words[i] == words[i + 1])
return repetitions / (len(words) - 1)

@staticmethod
def average_word_length(text: str) → float:
    """Average length of words in characters"""
words = text.split()
if len(words) == 0:
    return 0.0
return np.mean([len(w) for w in words])

@staticmethod
def sentence_count(text: str) → int:
    """Count sentences (periods, exclamation, question marks)"""
return text.count('!') + text.count('!') + text.count('?')

@staticmethod
def average_sentence_length(text: str) → float:
    """Average words per sentence"""
sentences = text.split('.')

```

```

words_per_sentence = [len(s.split()) for s in sentences if s.strip()]

if len(words_per_sentence) == 0:
    return 0.0

return np.mean(words_per_sentence)

@staticmethod
def vocabulary_size(text: str) → int:
    """Number of unique words"""
    return len(set(text.lower().split()))

# =====
# ADVANCED METRICS
# =====

class AdvancedMetrics:
    """More sophisticated metrics for generation quality"""

    @staticmethod
    def perplexity_from_loss(loss: float) → float:
        """
        Convert cross-entropy loss to perplexity

        Perplexity = e^loss
        """

        Args:
            loss: Cross-entropy loss value

        Returns:
            Perplexity score
        """

        return math.exp(loss)

    @staticmethod
    def self_bleu_score(generated_list: List[str], n_gram: int = 1) → float:
        """
        Compute SELF-BLEU (diversity among multiple generations)

        Measures: How different are multiple outputs from same seed?
        Lower = more diverse (good for creativity)
        Higher = less diverse (indicates overfitting)

        Args:
    
```

```

generated_list: List of generated texts from same seed
n_gram: N-gram size (1=unigrams, 2=bigrams)

Returns:
    SELF-BLEU score (0-1, lower is better)
"""

if len(generated_list) < 2:
    return 0.0

def get_ngrams(text, n):
    words = text.lower().split()
    return set(tuple(words[i:i+n]) for i in range(len(words) - n + 1))

# Compare each pair
similarities = []
for i in range(len(generated_list)):
    ngrams_i = get_ngrams(generated_list[i], n_gram)

    for j in range(i + 1, len(generated_list)):
        ngrams_j = get_ngrams(generated_list[j], n_gram)

        if len(ngrams_i) == 0 or len(ngrams_j) == 0:
            similarity = 0.0
        else:
            intersection = len(ngrams_i & ngrams_j)
            union = len(ngrams_i | ngrams_j)
            similarity = intersection / union if union > 0 else 0.0

        similarities.append(similarity)

return np.mean(similarities) if similarities else 0.0

@staticmethod
def vocabulary_coverage(generated_text: str, reference_text: str) → float:
"""

Compute: What fraction of unique reference words appear in generated text?

Range: 0-1
Higher = model uses similar vocabulary

Args:
    generated_text: Model output
    reference_text: Training data or reference

Returns:
    Coverage score
"""

```

```

gen_vocab = set(generated_text.lower().split())
ref_vocab = set(reference_text.lower().split())

if len(ref_vocab) == 0:
    return 0.0

overlap = len(gen_vocab & ref_vocab)
return overlap / len(ref_vocab)

@staticmethod
def entropy_score(text: str) → float:
    """
    Compute entropy of word distribution

    Range: 0 to infinity
    Higher = more diverse (good for generation)
    Lower = repetitive (bad)

    Args:
        text: Generated text

    Returns:
        Entropy score
    """
    words = text.lower().split()
    if len(words) == 0:
        return 0.0

    word_freq = Counter(words)
    probabilities = np.array(list(word_freq.values())) / len(words)

    entropy = -np.sum(probabilities * np.log(probabilities))
    return entropy

@staticmethod
def type_token_ratio_sliding(text: str, window: int = 50) → float:
    """
    Compute average TTR over sliding windows

    Measures: Vocabulary richness throughout text

    Args:
        text: Generated text
        window: Words per window

    Returns:
        Average TTR
    """

```

```

"""
words = text.lower().split()

if len(words) < window:
    return len(set(words)) / len(words)

ttrs = []
for i in range(len(words) - window + 1):
    window_words = words[i:i+window]
    ttr = len(set(window_words)) / len(window_words)
    ttrs.append(ttr)

return np.mean(ttrs)

# =====
===
# COMPREHENSIVE EVALUATION PIPELINE
# =====
===

def evaluate_single_generation(
    text: str,
    seed: str = None,
    reference: str = None,
    temperature: float = 1.0
) → Dict:
    """
    Comprehensive evaluation of a single generated text
    """

Args:
    text: Generated text
    seed: Original seed text (optional)
    reference: Reference/training text (optional)
    temperature: Temperature used for generation

Returns:
    Dictionary of all metrics
"""

results = {
    'temperature': temperature,

    # Basic metrics
    'length_chars': len(text),
    'length_words': len(text.split()),
    'sentences': TextMetrics.sentence_count(text),
}

```

```

# Vocabulary metrics
'distinctness': TextMetrics.distinctness(text),
'vocab_size': TextMetrics.vocabulary_size(text),
'avg_word_length': TextMetrics.average_word_length(text),
'repetition_ratio': TextMetrics.repetition_ratio(text),

# Advanced metrics
'entropy': AdvancedMetrics.entropy_score(text),
'avg_sentence_length': TextMetrics.average_sentence_length(text),
'ttr_sliding': AdvancedMetrics.type_token_ratio_sliding(text),
}

# Optional metrics if reference provided
if reference:
    results['vocabulary_coverage'] = AdvancedMetrics.vocabulary_coverage(text, reference)

return results

# =====
===
# EXAMPLE USAGE
# =====
===

if __name__ == "__main__":
    # Example texts
    text_low_quality = "the greatest difficulty.the other to maintain themselves."
    text_medium_quality = "the greatest becoming poor at any private person. he kept having followed a bandon of the army in susperition than recognio, having been seen also assist the venetians, and to rem ain into italy,[1] this city o"
    text_high_quality = "the great country who we fuit divides him, so that i any easily eled by bind in the beginning was cansterfor thy battle, he remains at your discourable."

    print("EXAMPLE EVALUATION:\n")

    # Evaluate each
    for text, label in [
        (text_low_quality, "Low Quality (T=0.5)"),
        (text_medium_quality, "Medium Quality (T=1.0)"),
        (text_high_quality, "High Quality (trained model)"),
    ]:
        print(f"\n{label}")
        print(f"Text: {text[:60]}...")
        print(f"-" * 70)

```

```

metrics = evaluate_single_generation(text, temperature=[0.5, 1.0, 1.5][[0, 1, 2][0]])

print(f" Length: {metrics['length_words']} words")
print(f" Distinctness: {metrics['distinctness']:.3f} (higher=better)")
print(f" Repetition: {metrics['repetition_ratio']:.3f} (lower=better)")
print(f" Entropy: {metrics['entropy']:.3f} (higher=more diverse)")

print("\n" + "="*70)
print("Use these metrics to evaluate your LSTM outputs!")
print("="*70)

```

- Evaluation output

EXAMPLE EVALUATION:

Low Quality (T=0.5)

Text: the greatest difficulty.the other to maintain themselves....

—
Length: 7 words
Distinctness: 1.000 (higher=better)
Repetition: 0.000 (lower=better)
Entropy: 1.946 (higher=more diverse)

Medium Quality (T=1.0)

Text: the greatest becoming poor at any private person. he kept ha...

—
Length: 36 words
Distinctness: 0.917 (higher=better)
Repetition: 0.000 (lower=better)
Entropy: 3.453 (higher=more diverse)

High Quality (trained model)

Text: the great country who we fuit divides him, so that i any eas...

—
Length: 28 words
Distinctness: 0.964 (higher=better)
Repetition: 0.000 (lower=better)
Entropy: 3.283 (higher=more diverse)

=====

Use these metrics to evaluate your LSTM outputs!

=====